Assignment #4

Task #1 – Run web servers behind ALB

1. Launch server one

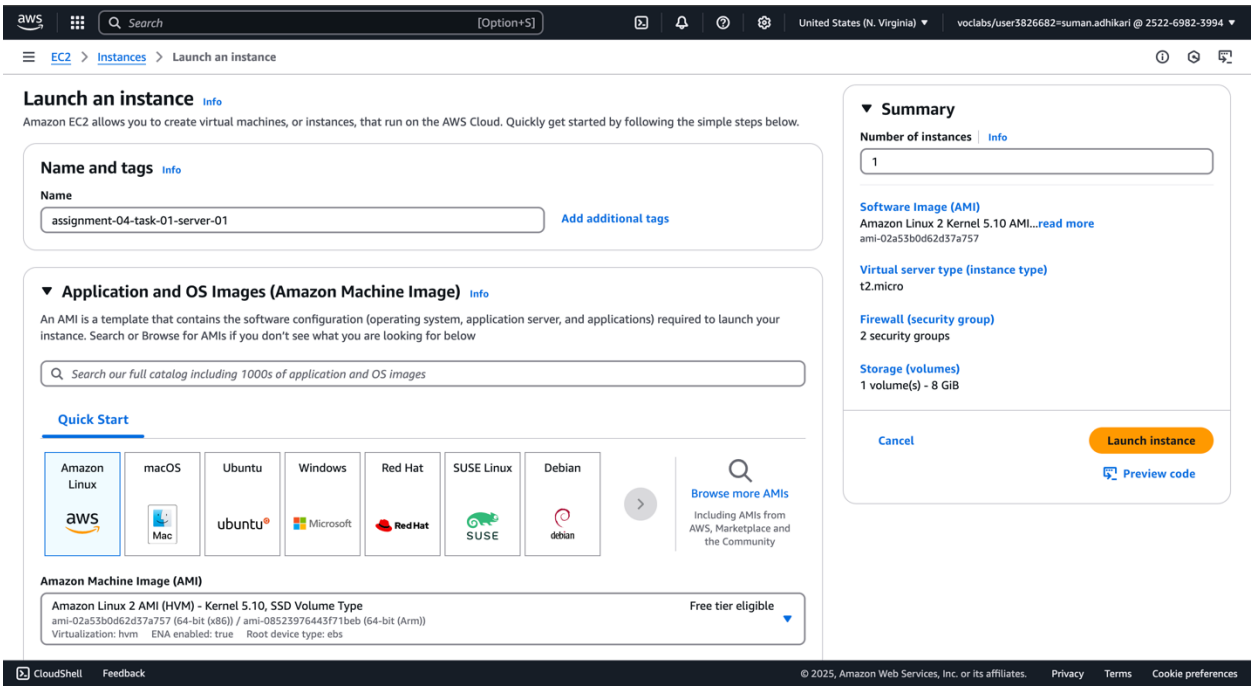

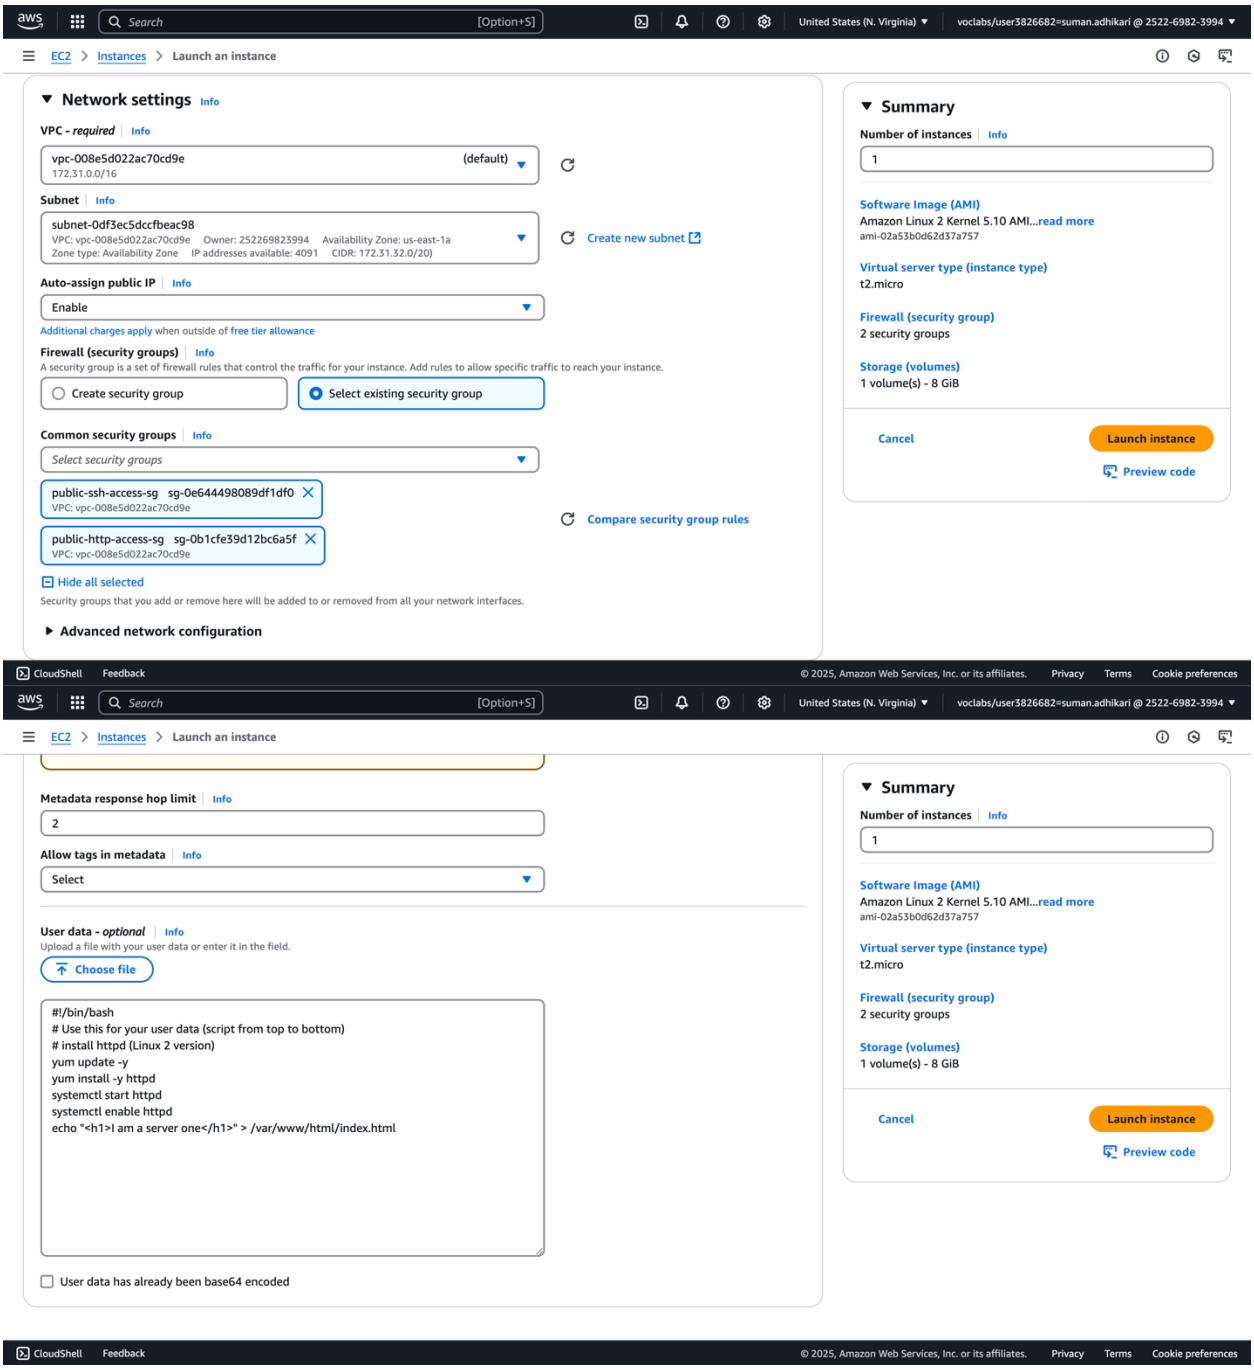
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Instances > Launch an instance
Network settings Info
VPC - required Info
vpc-008e5d022ac70cd9e (default)
172.31.0.0/16
Subnet Info
subnet-0df3ec5dccfbeac98
VPC: vpc-008e5d022ac70cd9e Owner: 252269823994 Availability Zone: us-east-1a
Zone type: Availability Zone IP addresses available: 4091 CIDR: 172.31.32.0/20)
Create new subnet
Auto-assign public IP Info
Enable
Additional charges apply when outside of free tier allowance
Firewall (security groups) Info
A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.
Create security group
Select existing security group
Common security groups Info
Select security groups
public-ssh-access-sg sg-0e644498089df1df0
VPC: vpc-008e5d022ac70cd9e
public-http-access-sg sg-0b1cfe39d12bc6a5f
VPC: vpc-008e5d022ac70cd9e
Compare security group rules
Hide all selected
Security groups that you add or remove here will be added to or removed from all your network interfaces.
Advanced network configuration
Summary
Number of instances Info
1
Software Image (AMI)
Amazon Linux 2 Kernel 5.10 AMI...read more
ami-02a53b0d62d37a757
Virtual server type (instance type)
t2.micro
Firewall (security group)
2 security groups
Storage (volumes)
1 volume(s) - 8 GiB
Cancel
Launch instance
Preview code
CloudShell Feedback
© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Instances > Launch an instance
Metadata response hop limit Info
2
Allow tags in metadata Info
Select
User data - optional Info
Upload a file with your user data or enter it in the field.
Choose file
#!/bin/bash
Use this for your user data (script from top to bottom)
install httpd (Linux 2 version)
yum update -y
yum install -y httpd
systemctl start httpd
systemctl enable httpd
echo "<h1>I am a server one</h1>" > /var/www/html/index.html
User data has already been base64 encoded
Summary
Number of instances Info
1
Software Image (AMI)
Amazon Linux 2 Kernel 5.10 AMI...read more
ami-02a53b0d62d37a757
Virtual server type (instance type)
t2.micro
Firewall (security group)
2 security groups
Storage (volumes)
1 volume(s) - 8 GiB
Cancel
Launch instance
Preview code
CloudShell Feedback
© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

2. Launch server two.

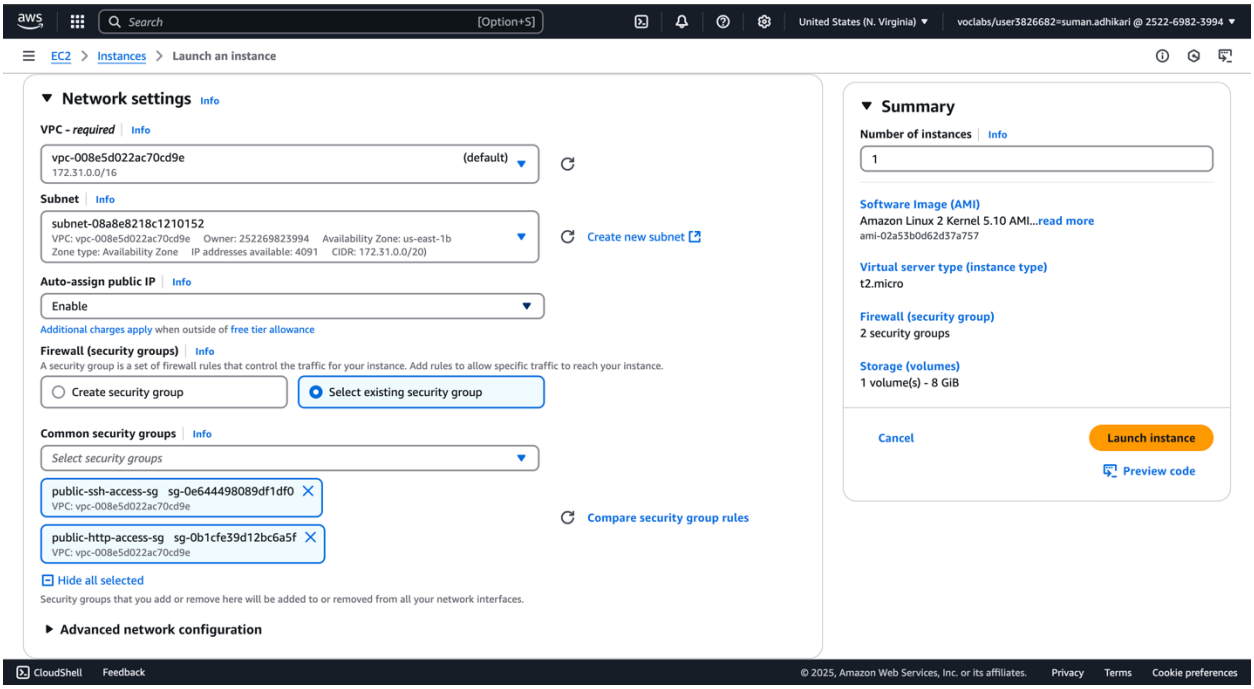

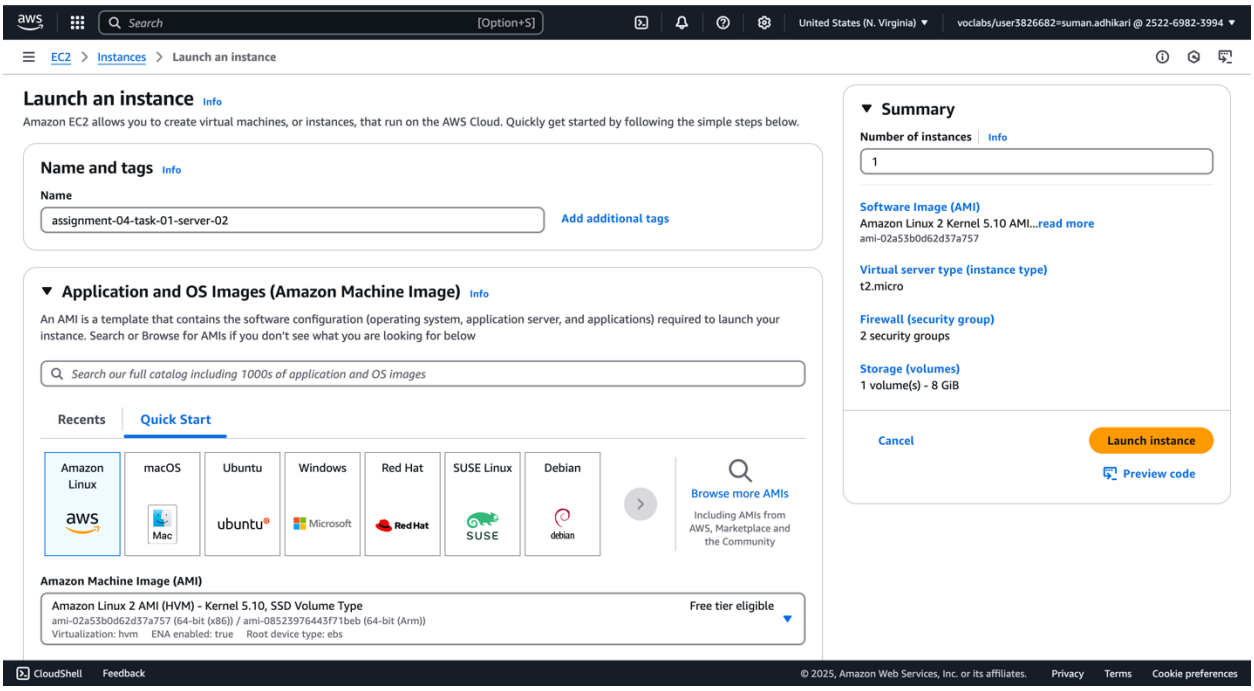

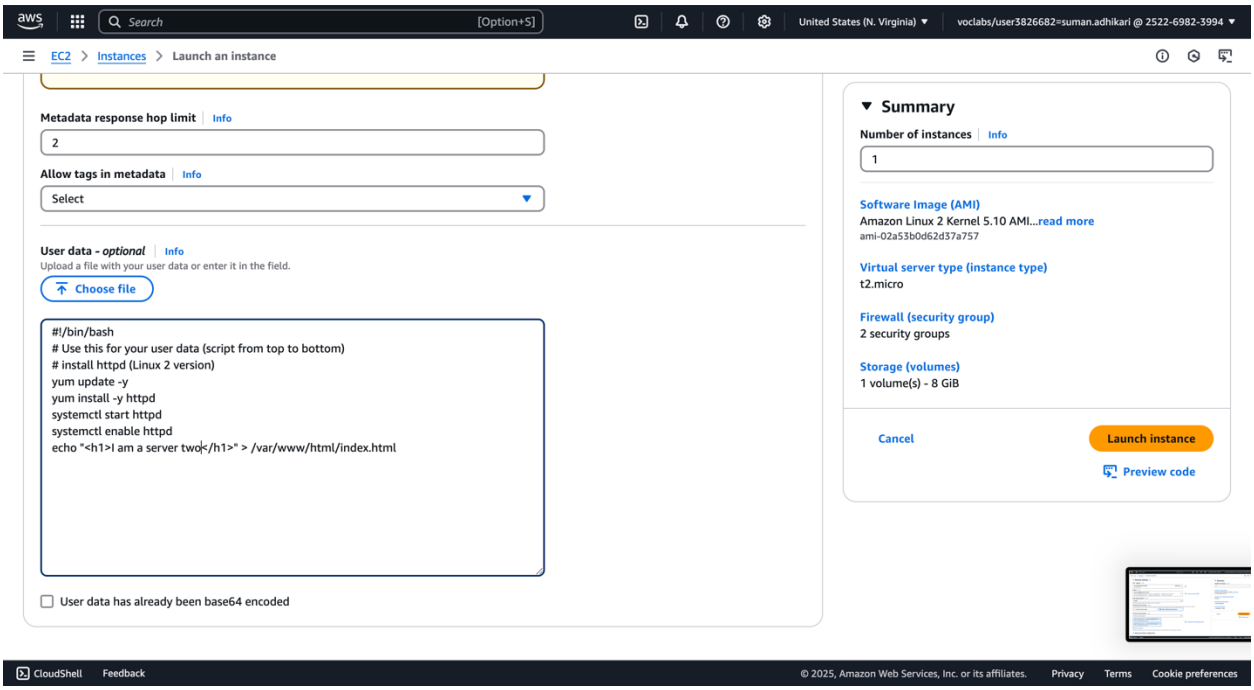

3. Check responses from two servers directly.

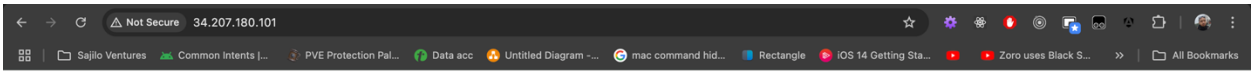

I am a server one

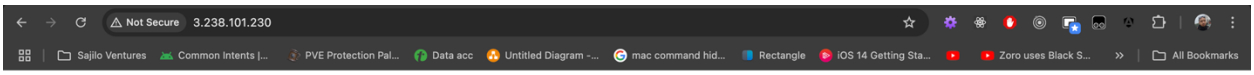

I am a server two

4. Create target group for ALB and register the two instances to the security groups. Add a health check on path / .

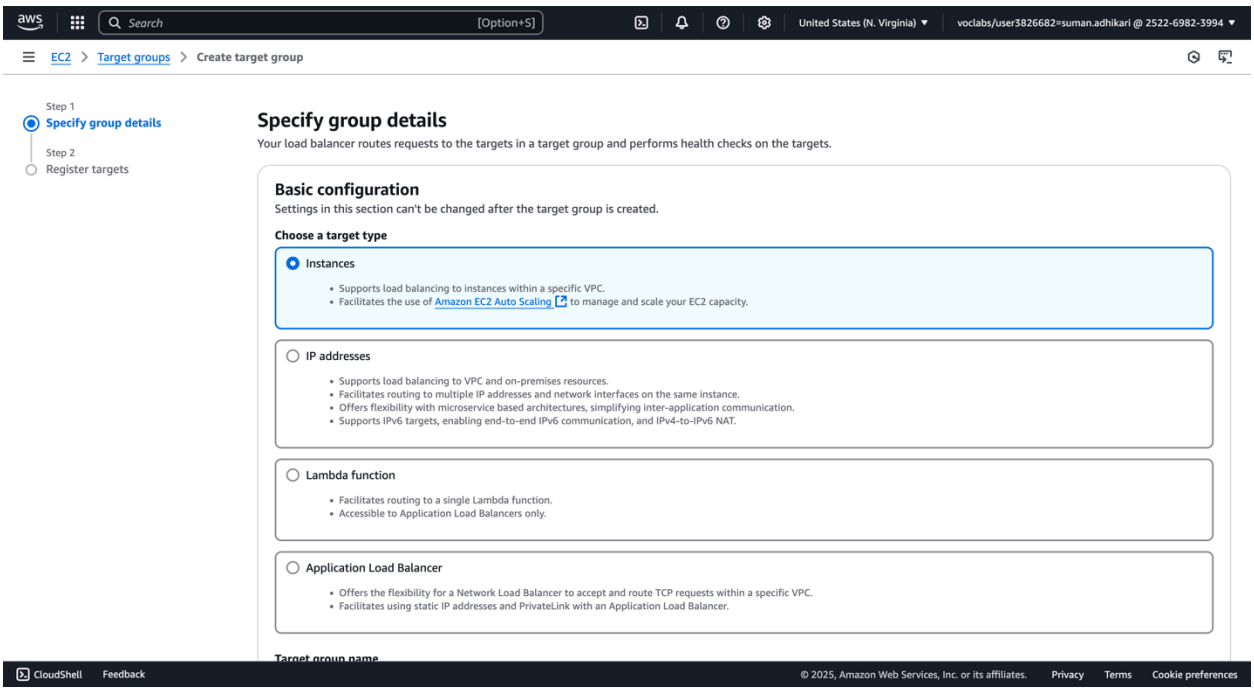

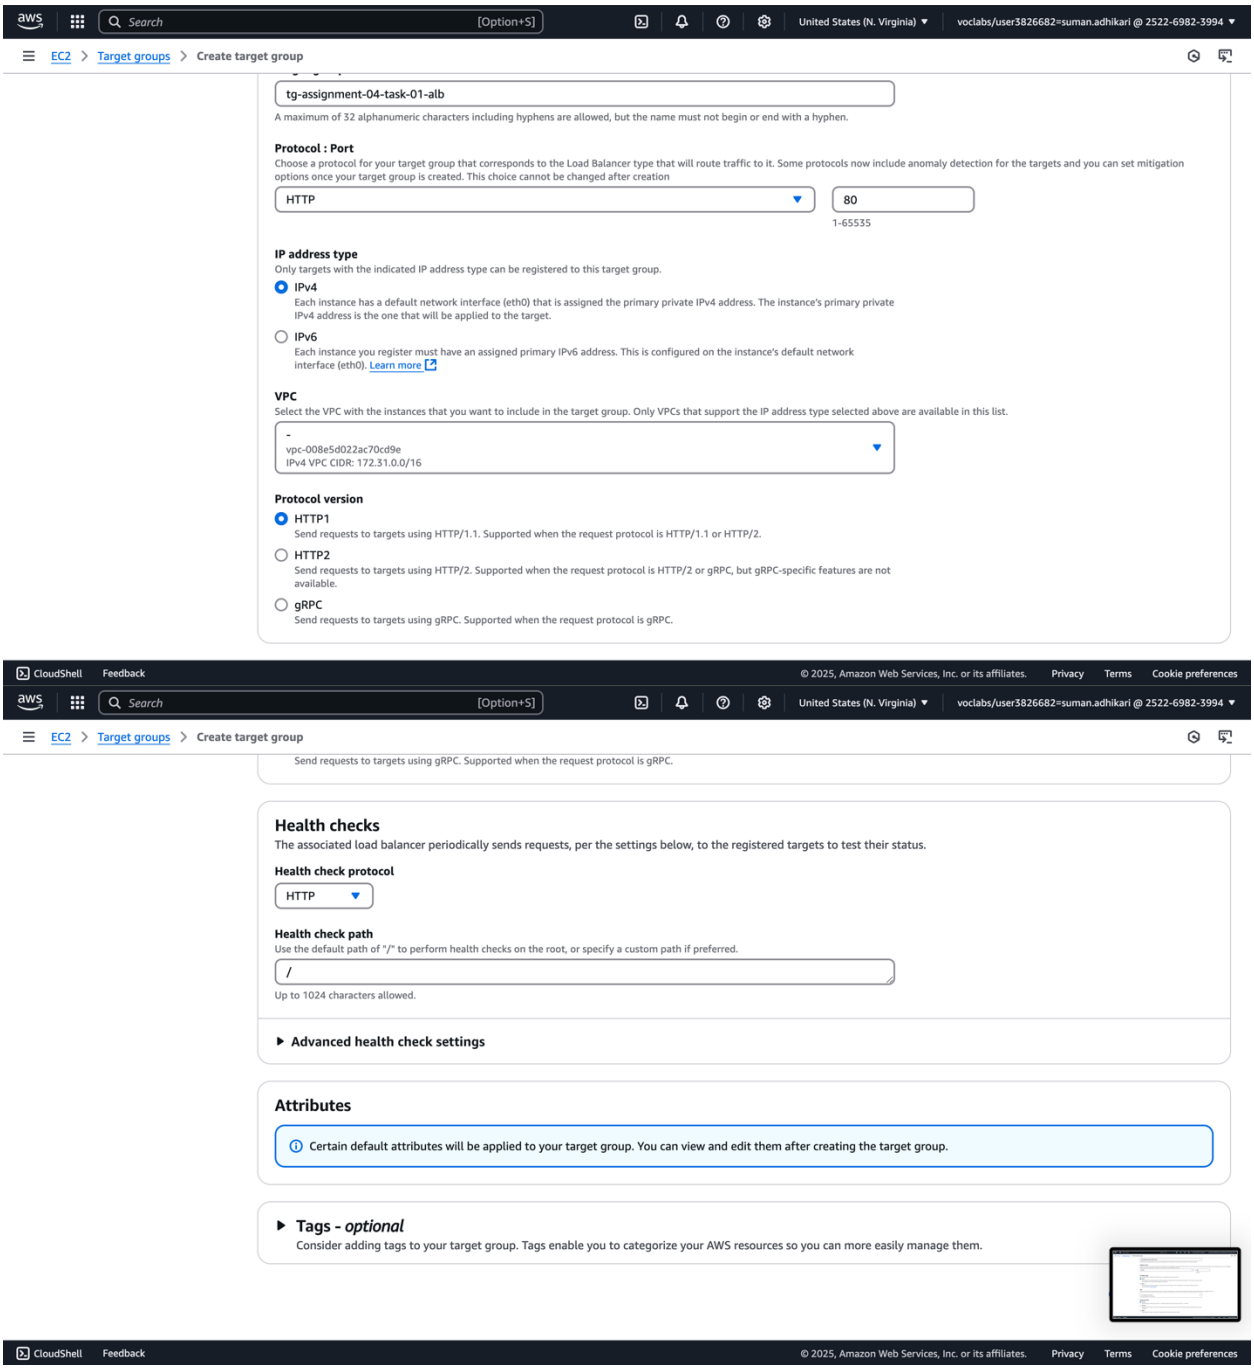
Search
[Option+5]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2
Target groups
Create target group
tg-assignment-04-task-01-alb
A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.
Protocol : Port
Choose a protocol for your target group that corresponds to the Load Balancer type that will route traffic to it. Some protocols now include anomaly detection for the targets and you can set mitigation options once your target group is created. This choice cannot be changed after creation
HTTP
80
1-65535
IP address type
Only targets with the indicated IP address type can be registered to this target group.
IPv4
Each instance has a default network interface (eth0) that is assigned the primary private IPv4 address. The instance's primary private IPv4 address is the one that will be applied to the target.
IPv6
Each instance you register must have an assigned primary IPv6 address. This is configured on the instance's default network interface (eth0). Learn more
VPC
Select the VPC with the instances that you want to include in the target group. Only VPCs that support the IP address type selected above are available in this list.
-
vpc-008e5d022ac70cd9e
IPv4 VPC CIDR: 172.31.0.0/16
Protocol version
HTTP1
Send requests to targets using HTTP/1.1. Supported when the request protocol is HTTP/1.1 or HTTP/2.
HTTP2
Send requests to targets using HTTP/2. Supported when the request protocol is HTTP/2 or gRPC, but gRPC-specific features are not available.
gRPC
Send requests to targets using gRPC. Supported when the request protocol is gRPC.
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

Search
[Option+5]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2
Target groups
Create target group
Send requests to targets using gRPC. Supported when the request protocol is gRPC.
Health checks
The associated load balancer periodically sends requests, per the settings below, to the registered targets to test their status.
Health check protocol
HTTP
Health check path
Use the default path of "/" to perform health checks on the root, or specify a custom path if preferred.
/
Up to 1024 characters allowed.
Advanced health check settings
Attributes
Certain default attributes will be applied to your target group. You can view and edit them after creating the target group.
Tags - optional
Consider adding tags to your target group. Tags enable you to categorize your AWS resources so you can more easily manage them.
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

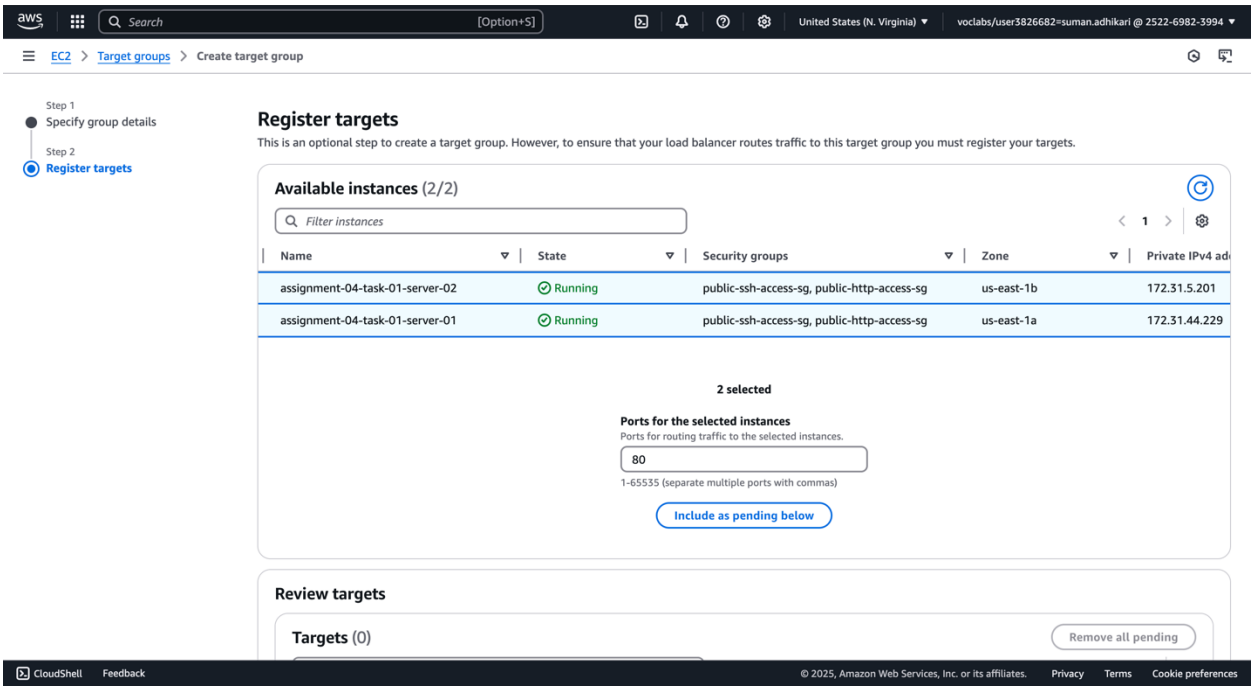

5. Target group is now created and in healthy state.

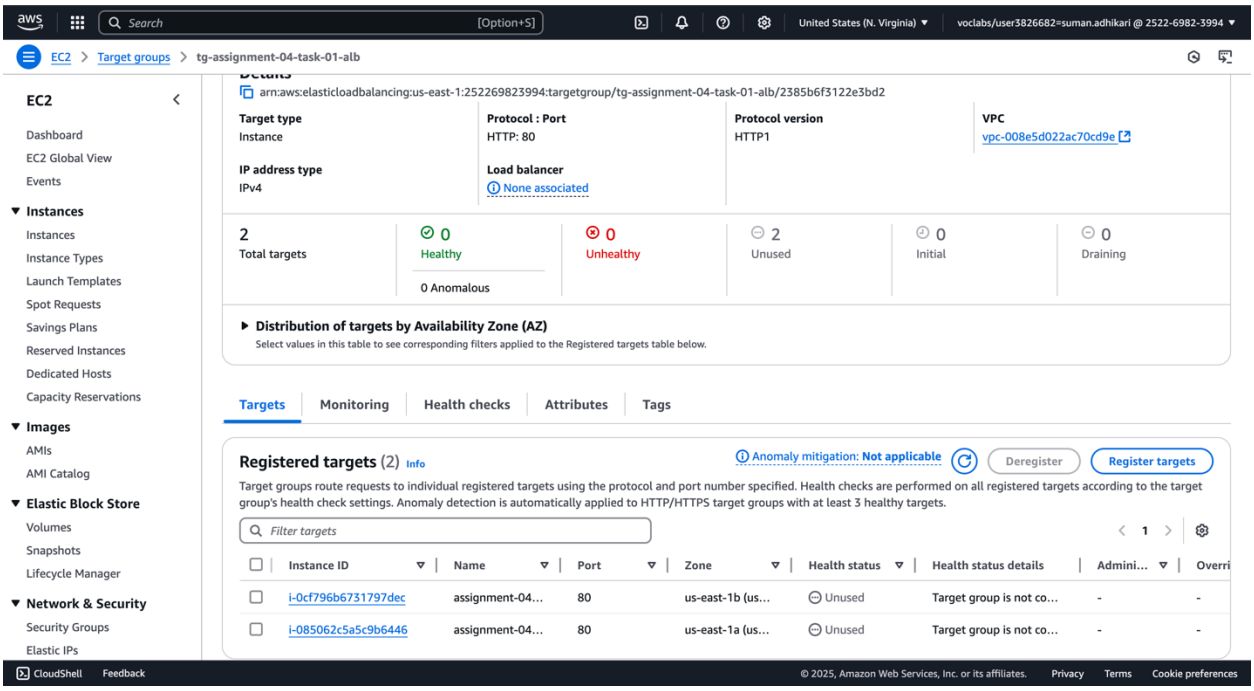

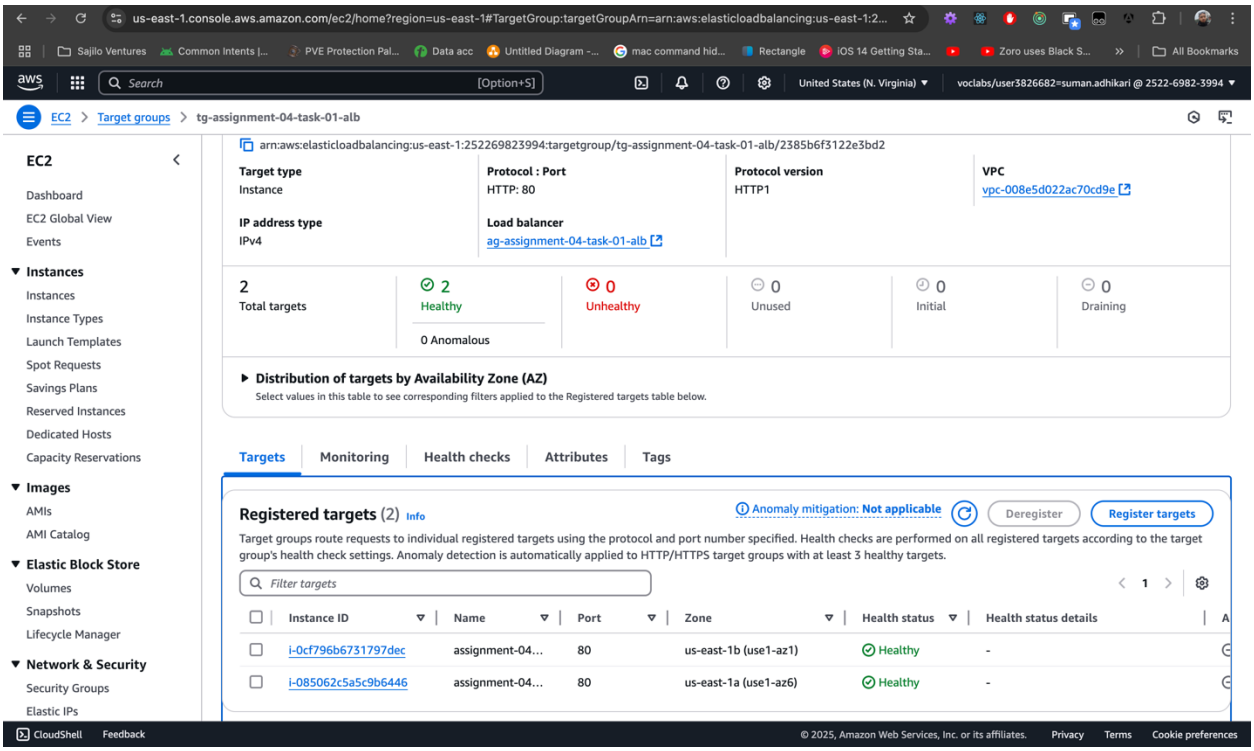

6. Create an Application Load Balancer.

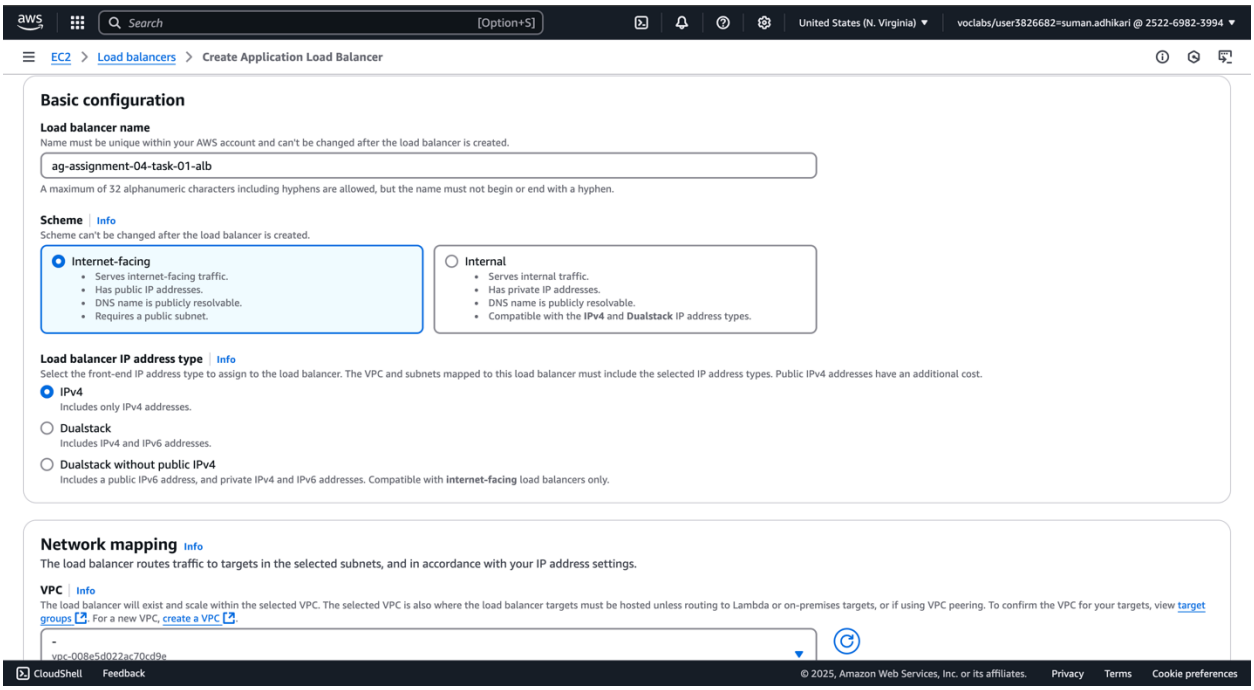

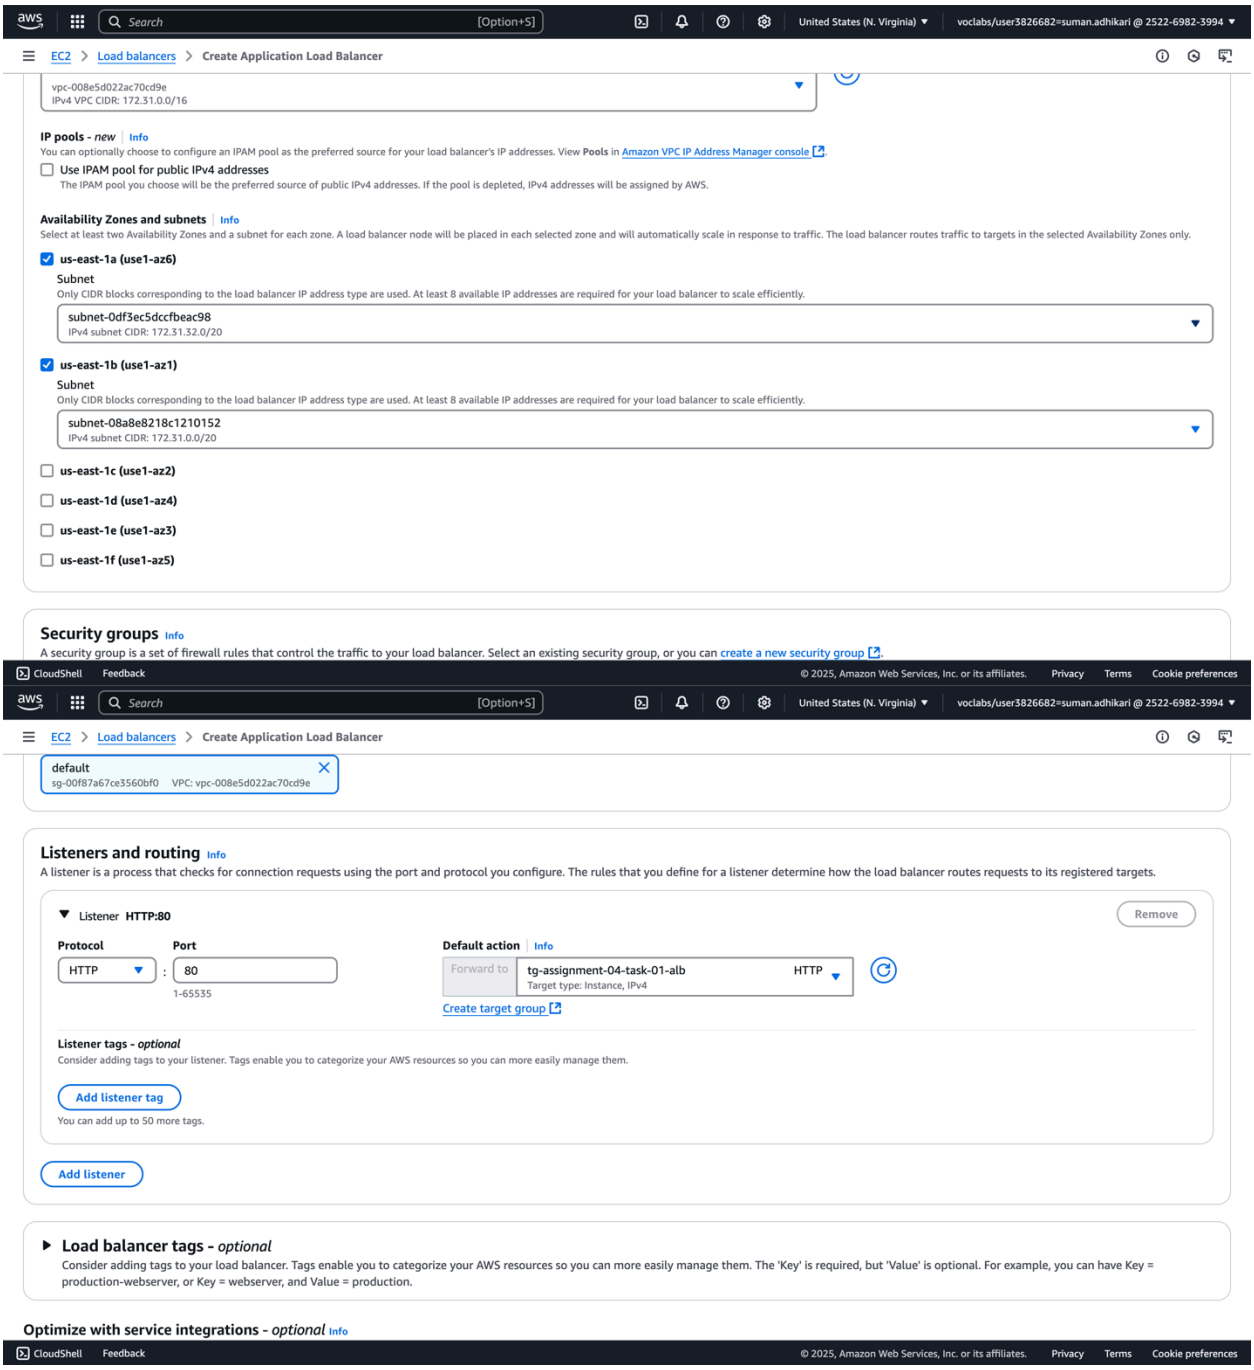
Search
[Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Load balancers > Create Application Load Balancer
vpc-008e5d022ac70cd9e
IPv4 VPC CIDR: 172.31.0.0/16
IP pools - new Info
You can optionally choose to configure an IPAM pool as the preferred source for your load balancer's IP addresses. View Pools in Amazon VPC IP Address Manager console.
Use IPAM pool for public IPv4 addresses
The IPAM pool you choose will be the preferred source of public IPv4 addresses. If the pool is depleted, IPv4 addresses will be assigned by AWS.
Availability Zones and subnets Info
Select at least two Availability Zones and a subnet for each zone. A load balancer node will be placed in each selected zone and will automatically scale in response to traffic. The load balancer routes traffic to targets in the selected Availability Zones only.
us-east-1a (use1-az6)
Subnet
Only CIDR blocks corresponding to the load balancer IP address type are used. At least 8 available IP addresses are required for your load balancer to scale efficiently.
subnet-0df3ec5dccfbeac98
IPv4 subnet CIDR: 172.31.32.0/20
us-east-1b (use1-az1)
Subnet
Only CIDR blocks corresponding to the load balancer IP address type are used. At least 8 available IP addresses are required for your load balancer to scale efficiently.
subnet-08a8e8218c1210152
IPv4 subnet CIDR: 172.31.0.0/20
us-east-1c (use1-az2)
us-east-1d (use1-az4)
us-east-1e (use1-az3)
us-east-1f (use1-az5)
Security groups Info
A security group is a set of firewall rules that control the traffic to your load balancer. Select an existing security group, or you can create a new security group.
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
default
sg-00f87a67ce3560bf0 VPC: vpc-008e5d022ac70cd9e
Listeners and routing Info
A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes requests to its registered targets.
Listener HTTP:80
Remove
Protocol
Port
HTTP
80
1-65535
Default action Info
Forward to
tg-assignment-04-task-01-alb
HTTP
Target type: Instance, IPv4
Create target group
Listener tags - optional
Consider adding tags to your listener. Tags enable you to categorize your AWS resources so you can more easily manage them.
Add listener tag
You can add up to 50 more tags.
Add listener
Load balancer tags - optional
Consider adding tags to your load balancer. Tags enable you to categorize your AWS resources so you can more easily manage them. The 'Key' is required, but 'Value' is optional. For example, you can have Key = production-webserver, or Key = webserver, and Value = production.
Optimize with service integrations - optional Info
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

7. ALB is now created. We are receiving response from ALB endpoint as well.

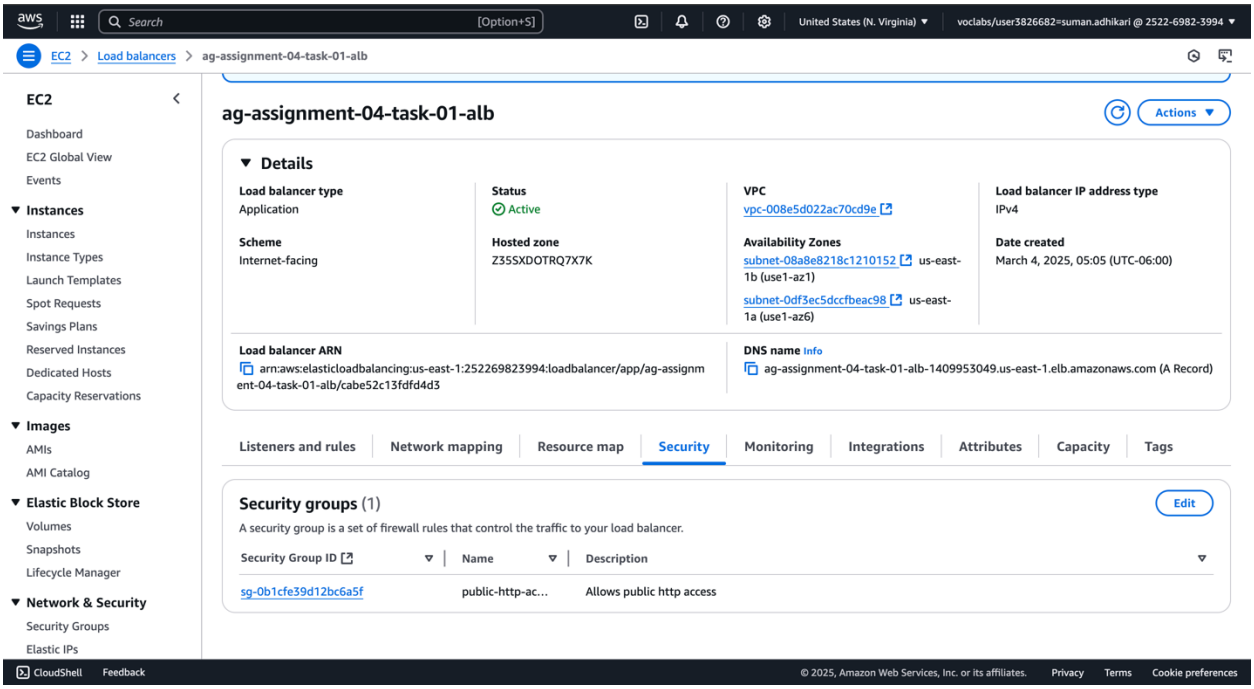

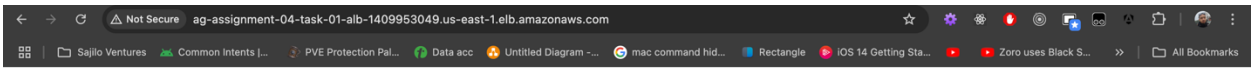

I am a server one

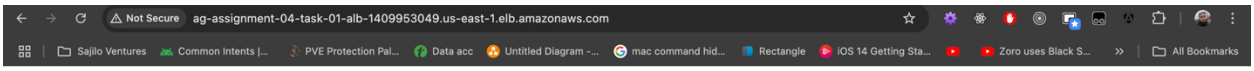

I am a server two

8. Now, hardening security to EC2 instances so that only ALB has access to it. Below we see updates to ALB security group and EC2 instances security group. ALB has ‘assignment-04-task-01-alb-sg’ group and EC2 instances have ‘private-assignment-04-task-01-alb-only-access’ group.

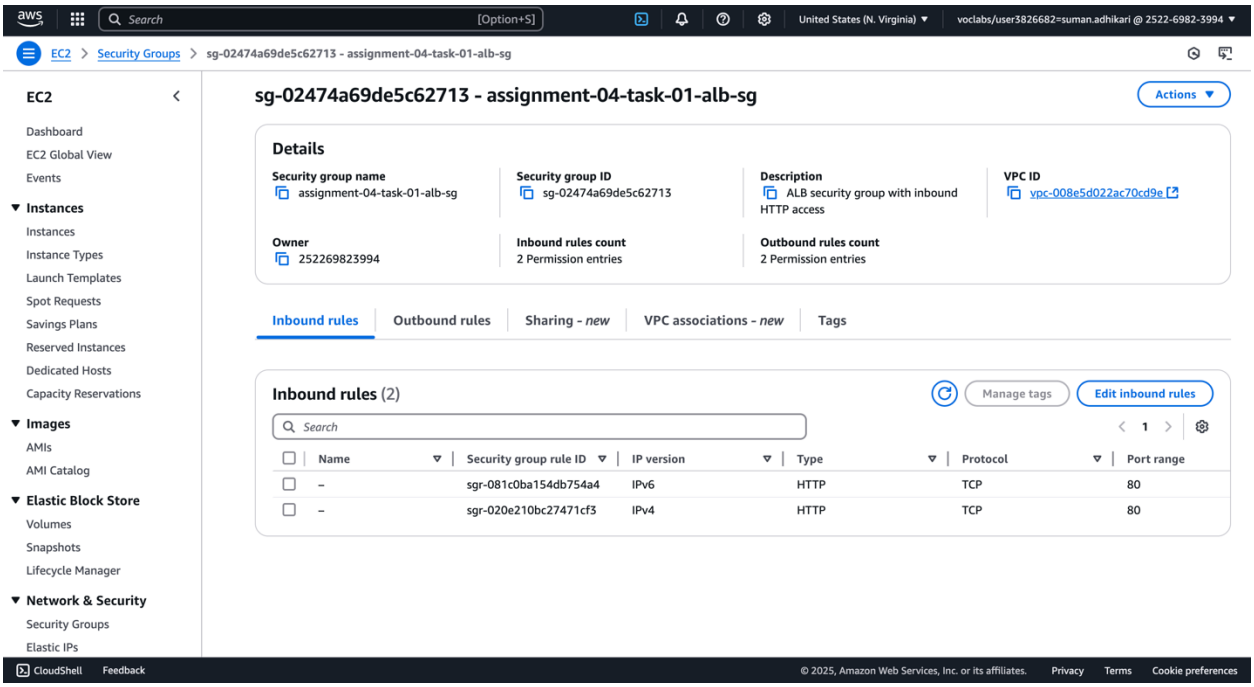

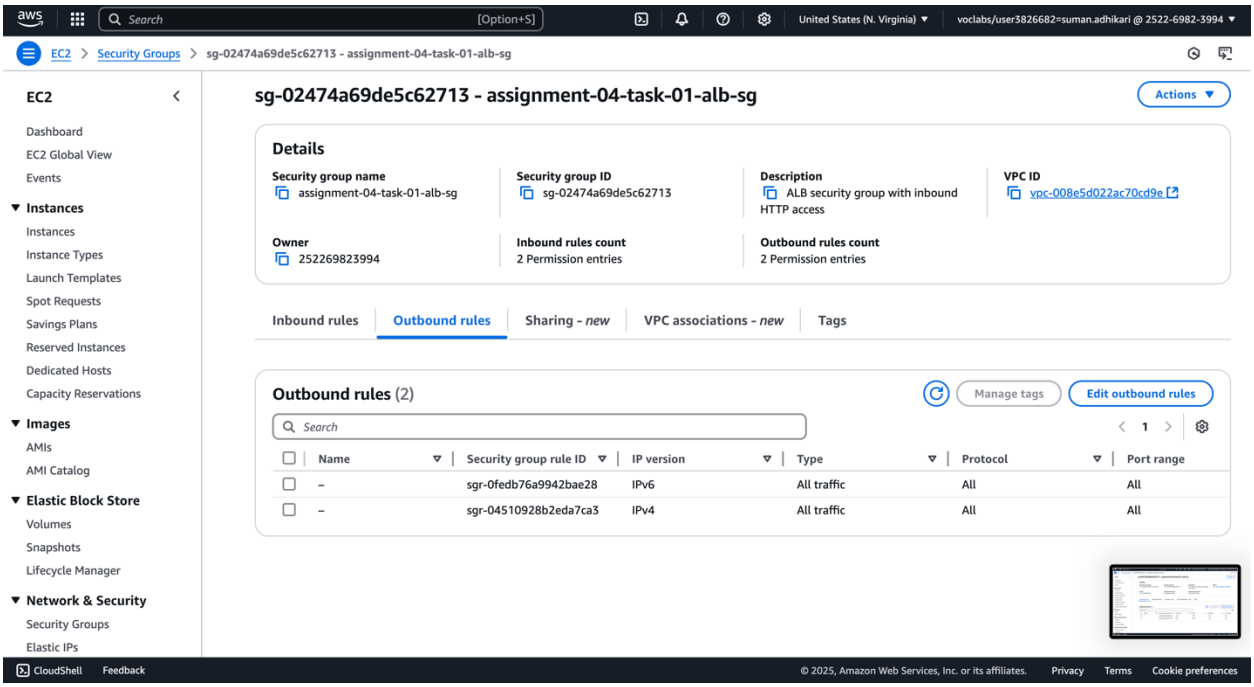
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Security Groups > sg-02474a69de5c62713 - assignment-04-task-01-alb-sg
EC2
Dashboard
EC2 Global View
Events
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
sg-02474a69de5c62713 - assignment-04-task-01-alb-sg
Actions
Details
Security group name
assignment-04-task-01-alb-sg
Security group ID
sg-02474a69de5c62713
Description
ALB security group with inbound HTTP access
VPC ID
vpc-008e5d022ac70cd9e
Owner
252269823994
Inbound rules count
2 Permission entries
Outbound rules count
2 Permission entries
Inbound rules
Outbound rules
Sharing - new
VPC associations - new
Tags
Outbound rules (2)
Manage tags
Edit outbound rules
Search
Name
Security group rule ID
IP version
Type
Protocol
Port range
– sgr-0fedb76a9942bae28 IPv6 All traffic All All
– sgr-04510928b2eda7ca3 IPv4 All traffic All All
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

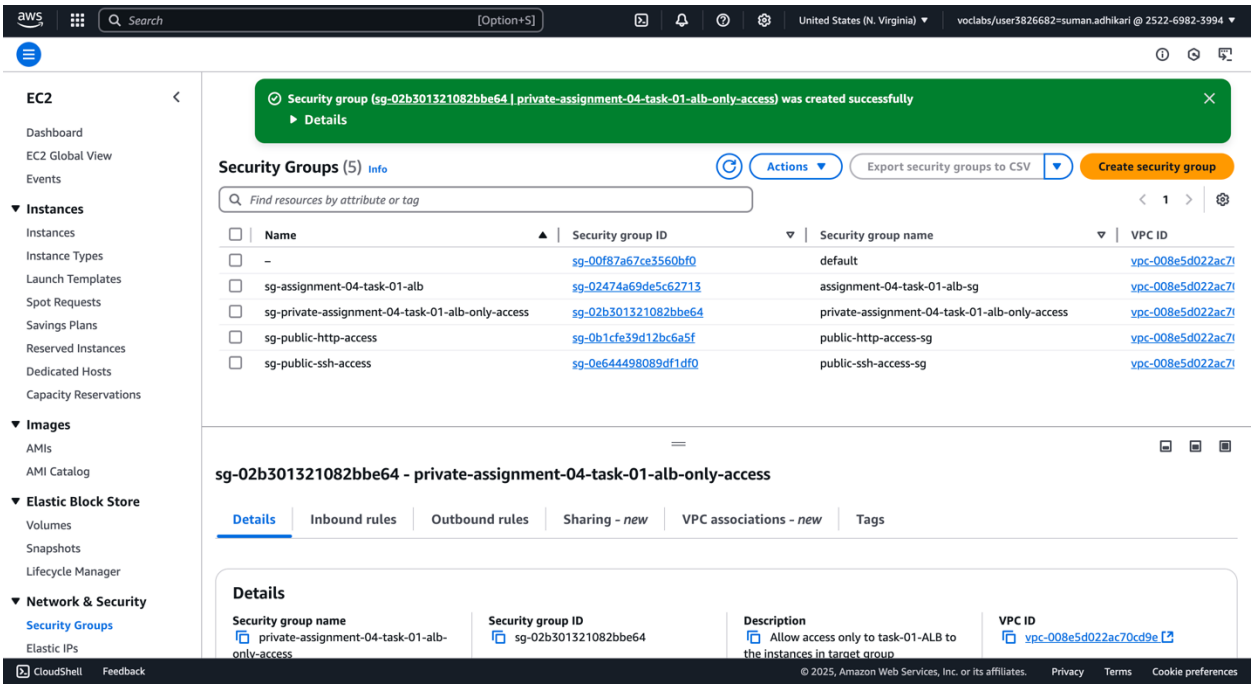
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2
Dashboard
EC2 Global View
Events
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Security group (sg-02b301321082bbe64 | private-assignment-04-task-01-alb-only-access) was created successfully
Details
Security Groups (5) Info
Actions
Export security groups to CSV
Create security group
Find resources by attribute or tag
Name
Security group ID
Security group name
VPC ID
– sg-00f87a67ce3560bf0 default vpc-008e5d022ac7
sg-assignment-04-task-01-alb sg-02474a69de5c62713 assignment-04-task-01-alb-sg vpc-008e5d022ac7
sg-private-assignment-04-task-01-alb-only-access sg-02b301321082bbe64 private-assignment-04-task-01-alb-only-access vpc-008e5d022ac7
sg-public-http-access sg-0b1cfe39d12bc6a5f public-http-access-sg vpc-008e5d022ac7
sg-public-ssh-access sg-0e644498089df1df0 public-ssh-access-sg vpc-008e5d022ac7
sg-02b301321082bbe64 - private-assignment-04-task-01-alb-only-access
Details
Inbound rules
Outbound rules
Sharing - new
VPC associations - new
Tags
Details
Security group name
private-assignment-04-task-01-alb-only-access
Security group ID
sg-02b301321082bbe64
Description
Allow access only to task-01-ALB to the instances in target group
VPC ID
vpc-008e5d022ac70cd9e
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

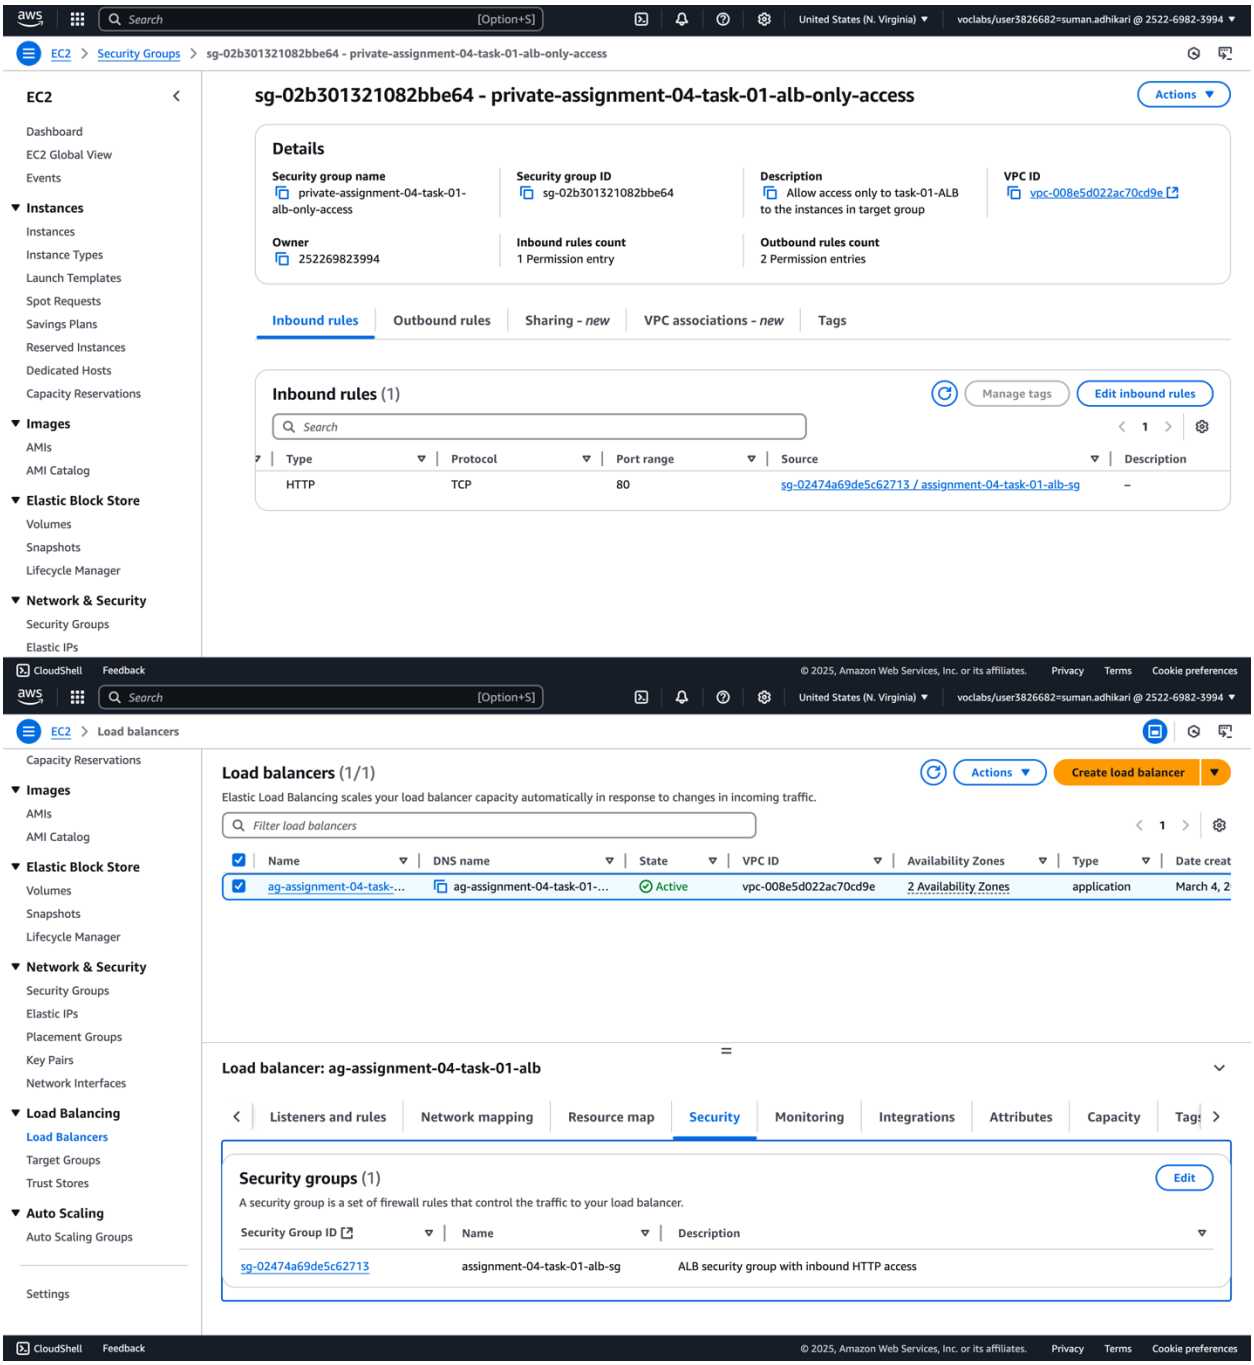
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Security Groups > sg-02b301321082bbe64 - private-assignment-04-task-01-alb-only-access
EC2
Dashboard
EC2 Global View
Events
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
sg-02b301321082bbe64 - private-assignment-04-task-01-alb-only-access
Actions
Details
Security group name
private-assignment-04-task-01-alb-only-access
Security group ID
sg-02b301321082bbe64
Description
Allow access only to task-01-ALB to the instances in target group
VPC ID
vpc-008e5d022ac70cd9e
Owner
252269823994
Inbound rules count
1 Permission entry
Outbound rules count
2 Permission entries
Inbound rules
Outbound rules
Sharing - new
VPC associations - new
Tags
Inbound rules (1)
Manage tags
Edit inbound rules
Search
Type
Protocol
Port range
Source
Description
HTTP
TCP
80
sg-02474a69de5c62713 / assignment-04-task-01-alb-sg
-
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
EC2 > Load balancers
Load balancers (1/1)
Actions
Create load balancer
Elastic Load Balancing scales your load balancer capacity automatically in response to changes in incoming traffic.
Filter load balancers
Name
DNS name
State
VPC ID
Availability Zones
Type
Date creat
ag-assignment-04-task-...
ag-assignment-04-task-01-...
Active
vpc-008e5d022ac70cd9e
2 Availability Zones
application
March 4, 2
Load balancer: ag-assignment-04-task-01-alb
Listeners and rules
Network mapping
Resource map
Security
Monitoring
Integrations
Attributes
Capacity
Security groups (1)
Edit
A security group is a set of firewall rules that control the traffic to your load balancer.
Security Group ID
Name
Description
sg-02474a69de5c62713
assignment-04-task-01-alb-sg
ALB security group with inbound HTTP access
Load Balancing
Load Balancers
Target Groups
Trust Stores
Auto Scaling
Auto Scaling Groups
Settings
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces

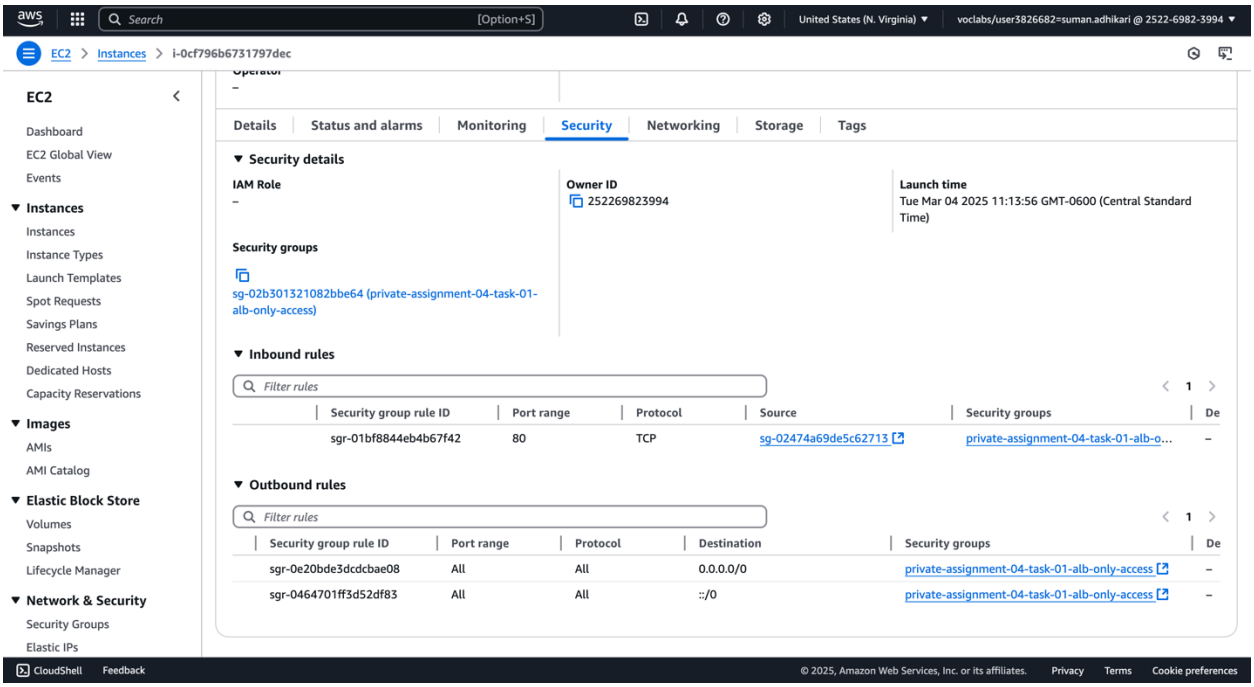

9. Check response from ALB now. We are receiving response only from ALB now.

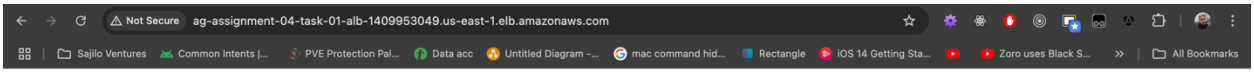

I am a server two

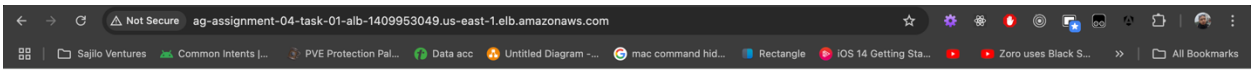

I am a server one

Task #2 – Run web servers behind NLB

We will be re-using the instances from Task #1. In summary, we will only create a new security group for EC2 instances to receive traffic from NLB and a security group for the NLB along with a new target group with TCP protocol for NLB to utilize. Everything else will work as Task #1 instances are already preconfigured with HTTP servers.

1. Create security group for NLB.

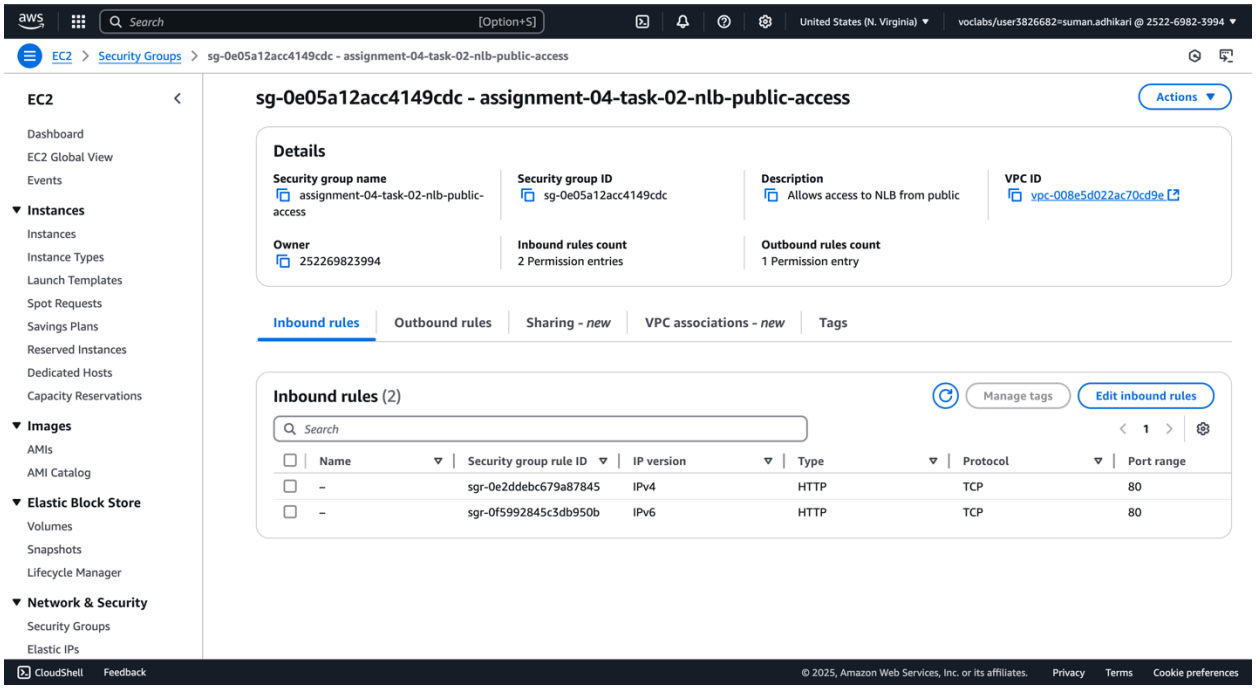

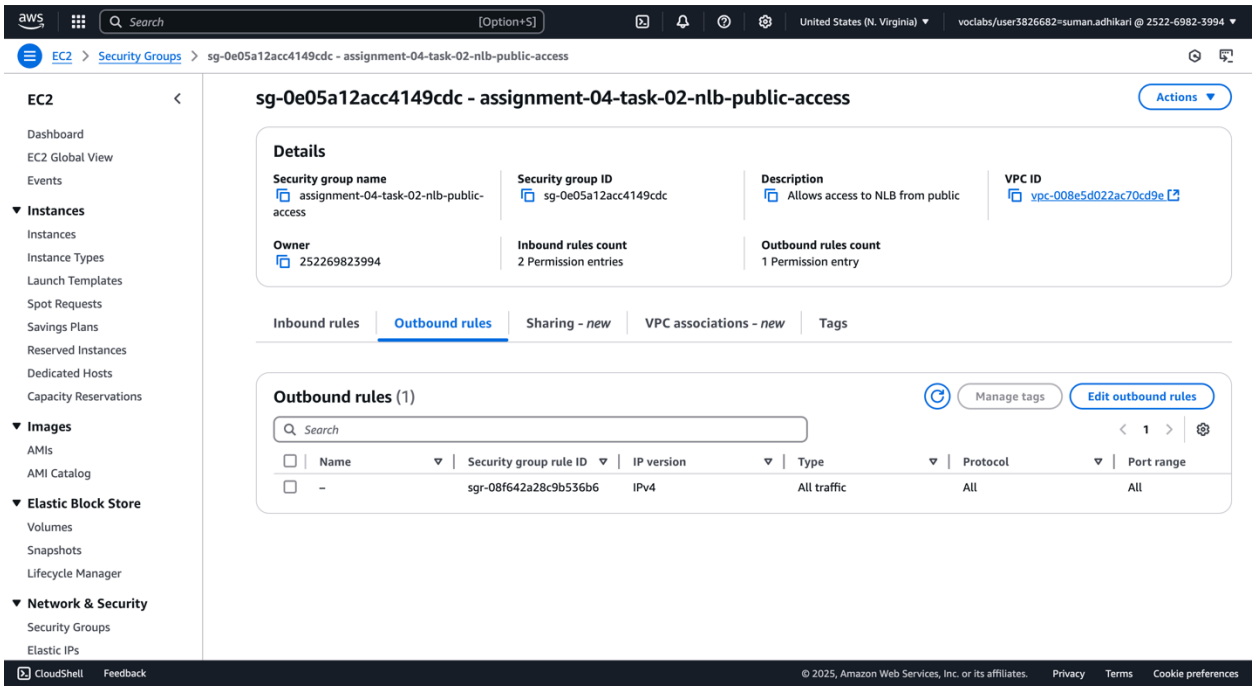

2. Create security group for EC2 instances to allow connection from NLB using TCP.

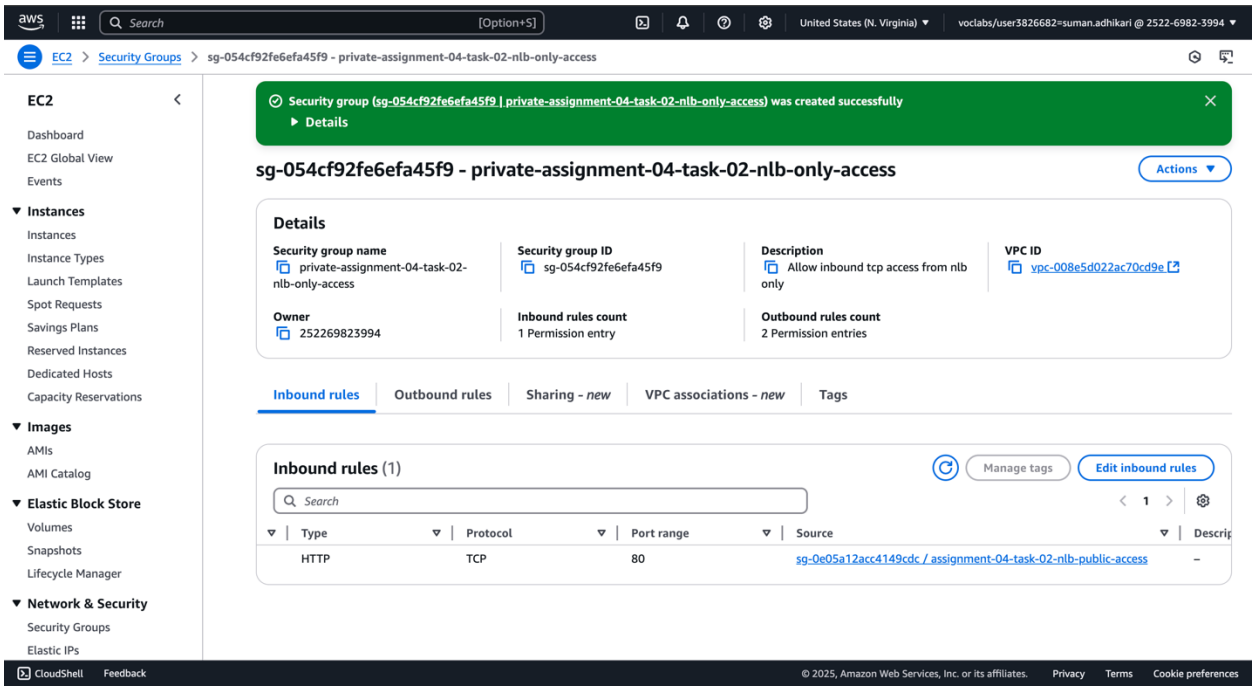

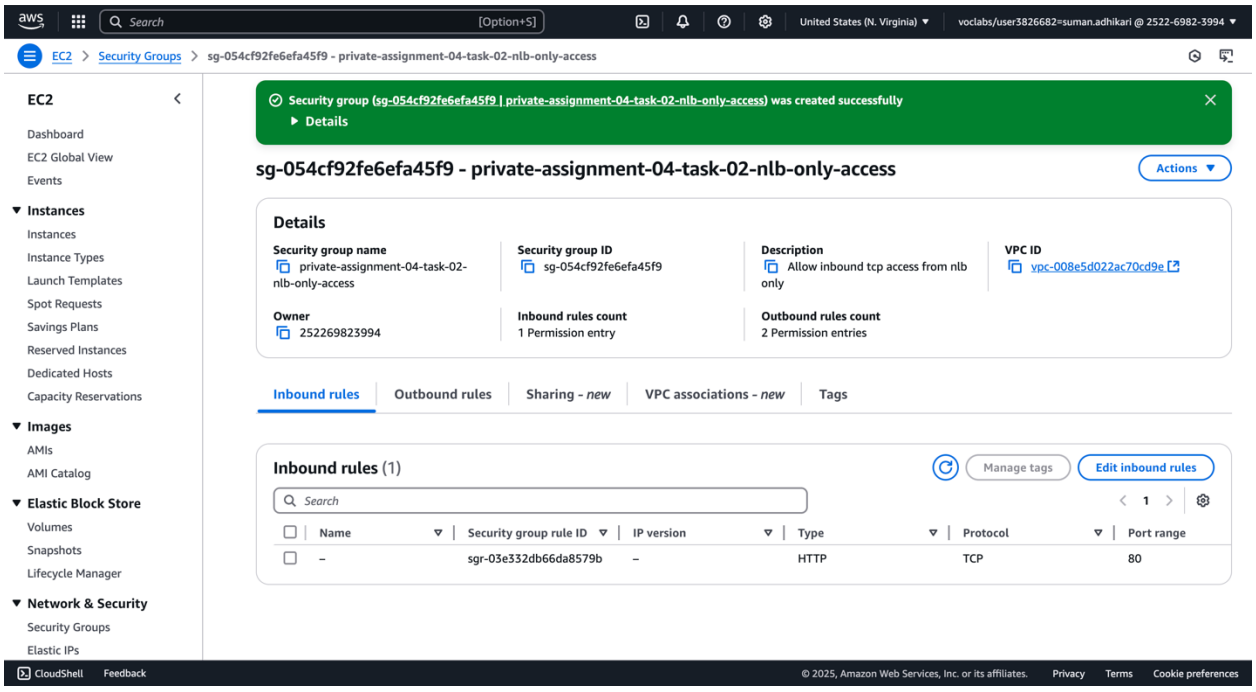

3. Attach ‘private-assignment-04-task-02-nlb-only-access’ security group to instances.

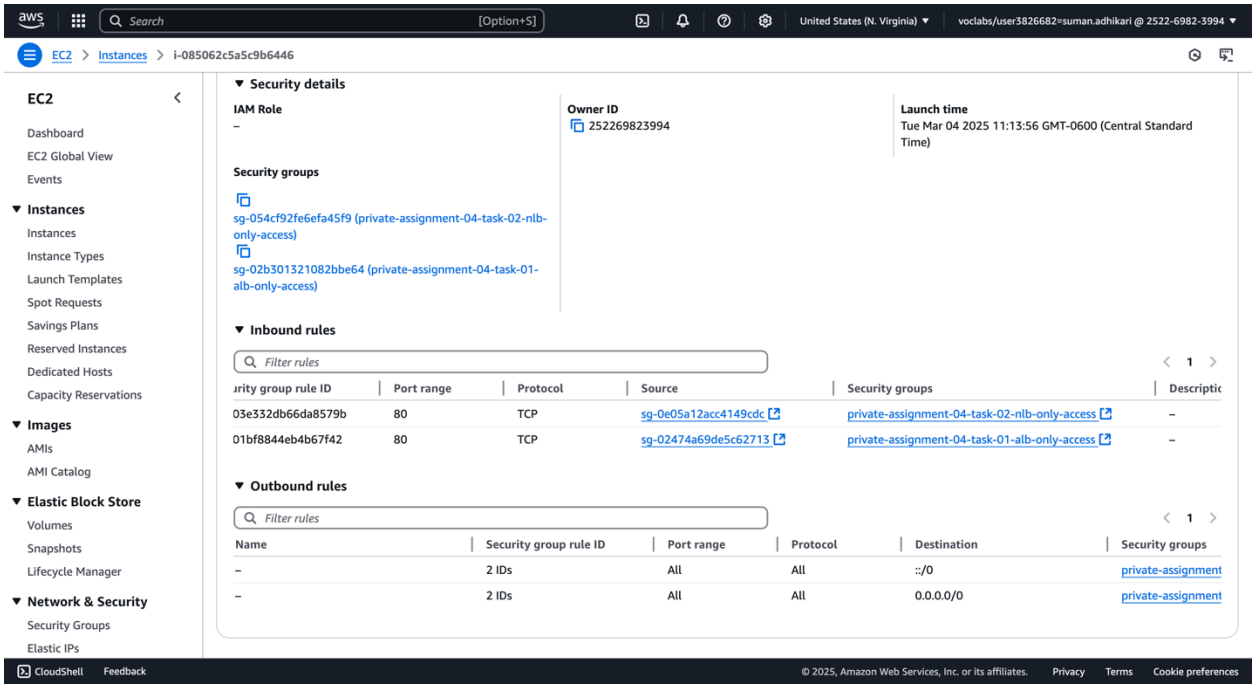

4. Create a target group for NLB and register instances with it.

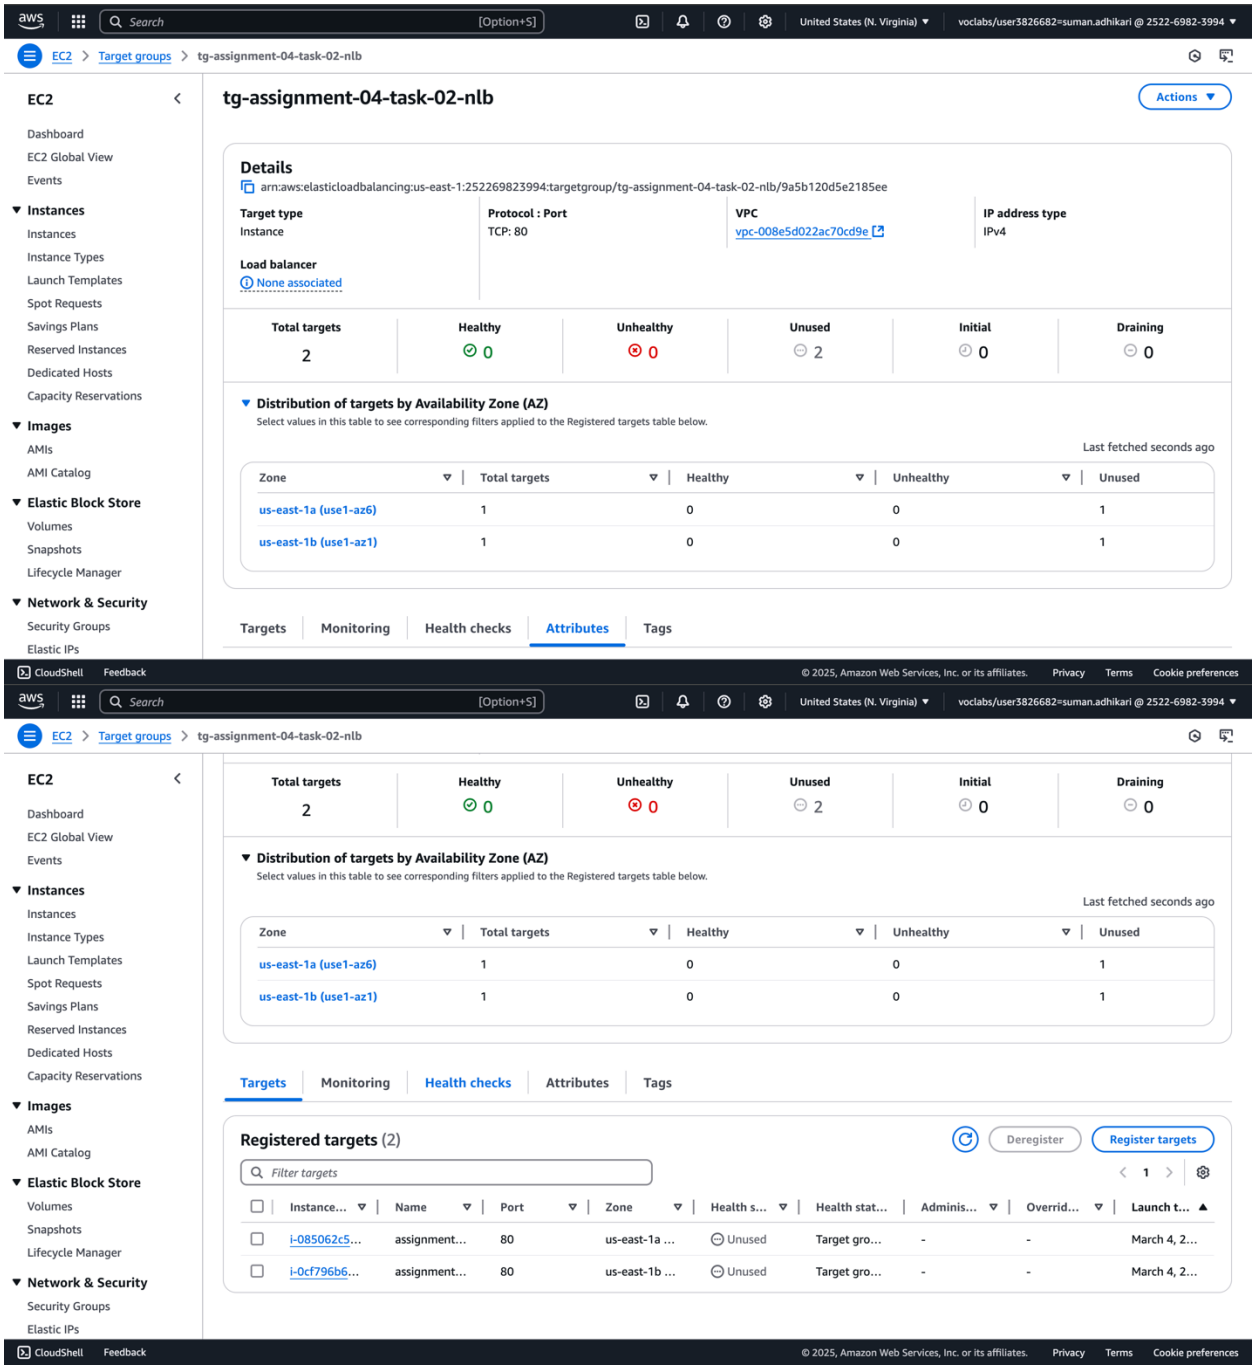

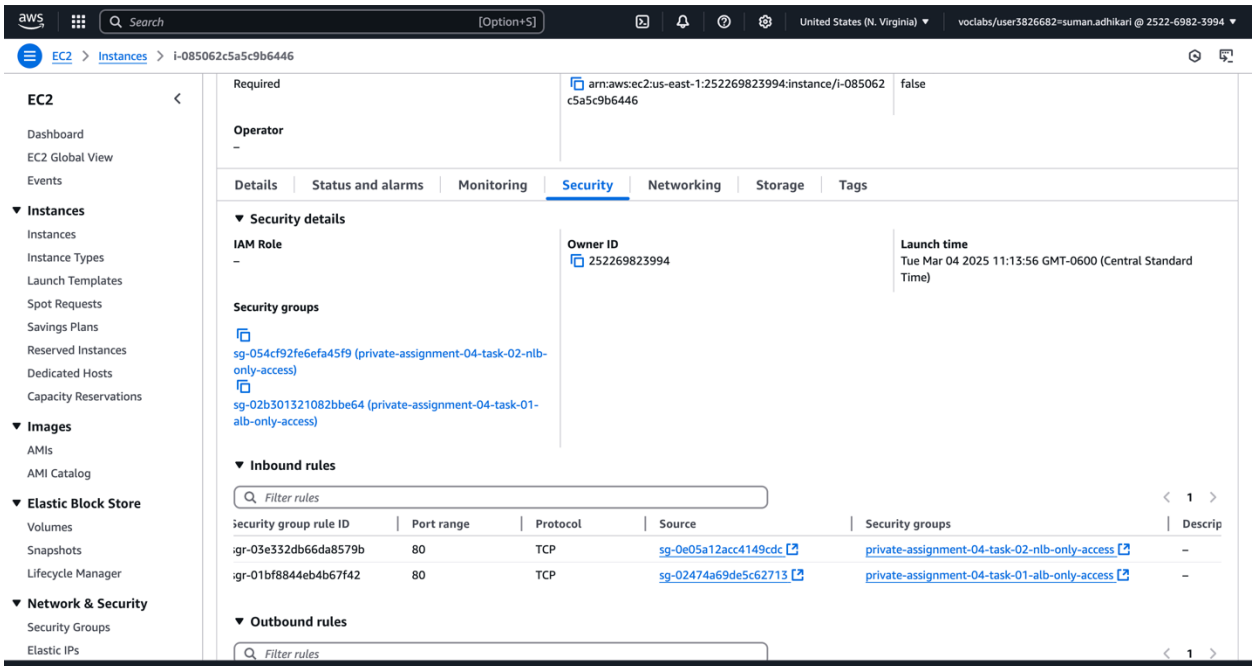
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Instances > i-085062c5a5c9b6446
EC2
Dashboard
EC2 Global View
Events
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Required
Operator
–
arn:aws:ec2:us-east-1:252269823994:instance/i-085062c5a5c9b6446
false
Details
Status and alarms
Monitoring
Security
Networking
Storage
Tags
Security details
IAM Role
–
Owner ID
252269823994
Launch time
Tue Mar 04 2025 11:13:56 GMT-0600 (Central Standard Time)
Security groups
sg-054cf92fe6efa45f9 (private-assignment-04-task-02-nlb-only-access)
sg-02b301321082bbe64 (private-assignment-04-task-01-alb-only-access)
Inbound rules
Filter rules
Security group rule ID
Port range
Protocol
Source
Security groups
Descrip
sgr-03e332db66da8579b
80
TCP
sg-0e05a12acc4149cdc
private-assignment-04-task-02-nlb-only-access
–
sgr-01bf8844eb4b67f42
80
TCP
sg-02474a69de5c62713
private-assignment-04-task-01-alb-only-access
–
Outbound rules
Filter rules
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

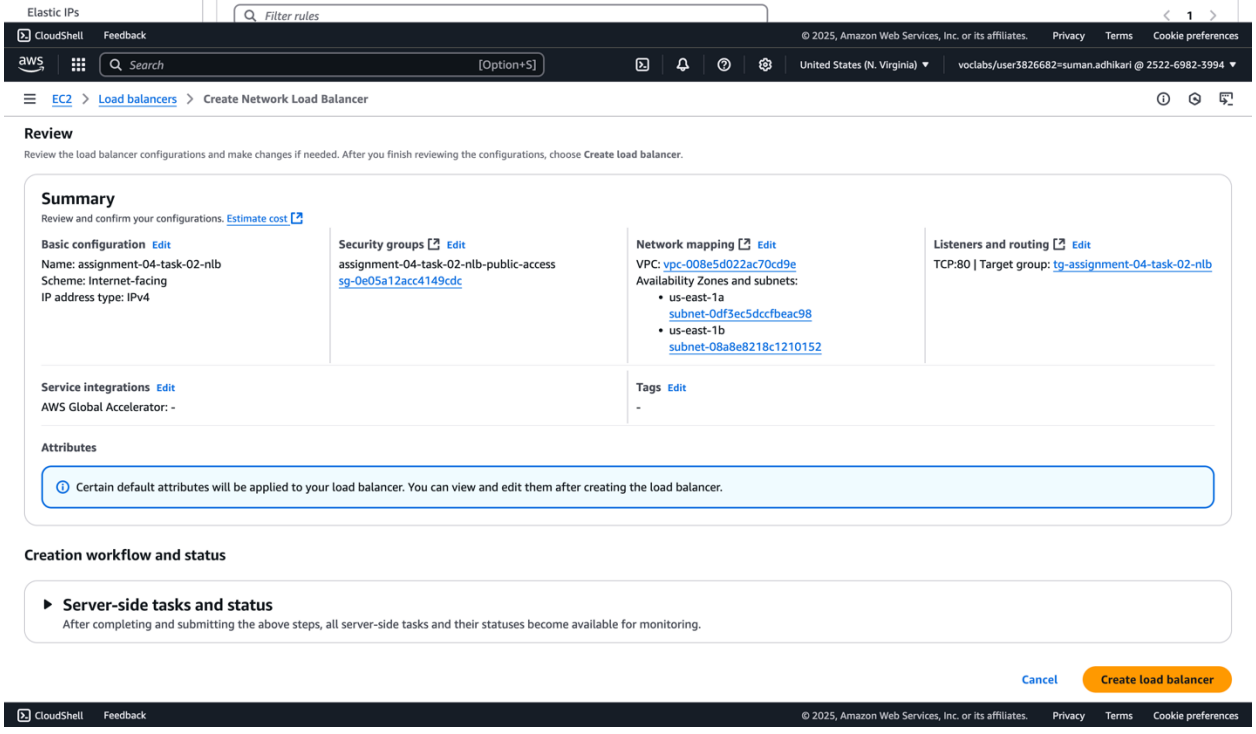
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Load balancers > Create Network Load Balancer
Review
Review the load balancer configurations and make changes if needed. After you finish reviewing the configurations, choose Create load balancer.
Summary
Review and confirm your configurations. Estimate cost
Basic configuration Edit
Name: assignment-04-task-02-nlb
Scheme: Internet-facing
IP address type: IPv4
Security groups Edit
assignment-04-task-02-nlb-public-access
sg-0e05a12acc4149cdc
Network mapping Edit
VPC: vpc-008e5d022ac70cd9e
Availability Zones and subnets:
us-east-1a
subnet-0df3ec5dccfbeac98
us-east-1b
subnet-08a8e8218c1210152
Listeners and routing Edit
TCP:80 | Target group: tg-assignment-04-task-02-nlb
Service integrations Edit
AWS Global Accelerator: -
Tags Edit
-
Attributes
Certain default attributes will be applied to your load balancer. You can view and edit them after creating the load balancer.
Creation workflow and status
Server-side tasks and status
After completing and submitting the above steps, all server-side tasks and their statuses become available for monitoring.
Cancel
Create load balancer
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

5. Instances are now healthy. And we are getting response from NLB.

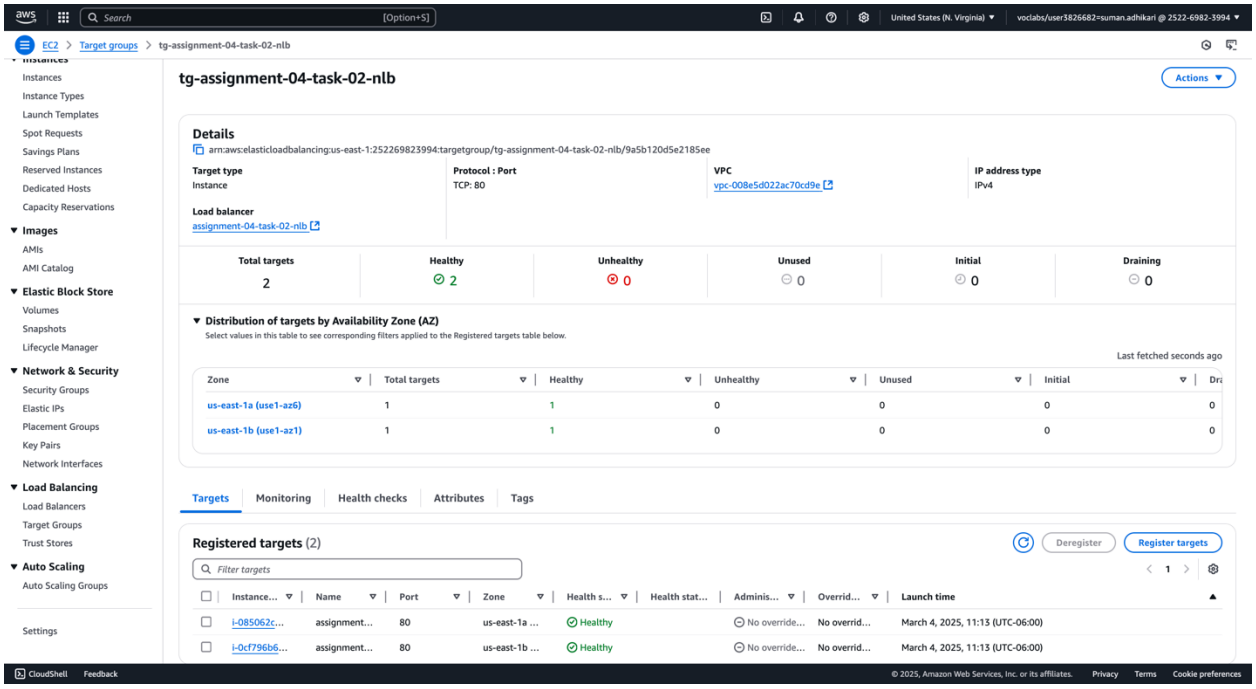

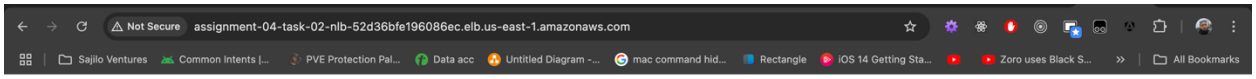

I am a server one

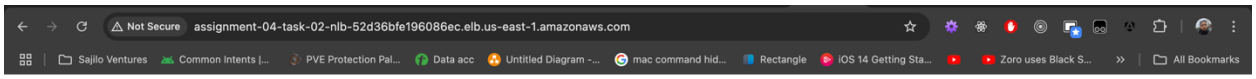

I am a server two

Task #3 – Run the web server in ASG

Here, we will first need a launch template for the ASG to run on. Then we will need to create the ASG itself and assign the launch template to it. While assigning the launch template, for creating ASG, we will need to select the VPC and availability zones and the target group we will use for the ASG to manage. Along with that we need to specify the desired, min and max instances to have and a target metric for Cloud Watch to check for to trigger the auto scaling. We will use 60% CPU utilization for that and will be using the stress tool to simulate the CPU load.

Note: we will be reusing the ALB and target group from task #1 for cost and time effectiveness.

1. Create a launch template.

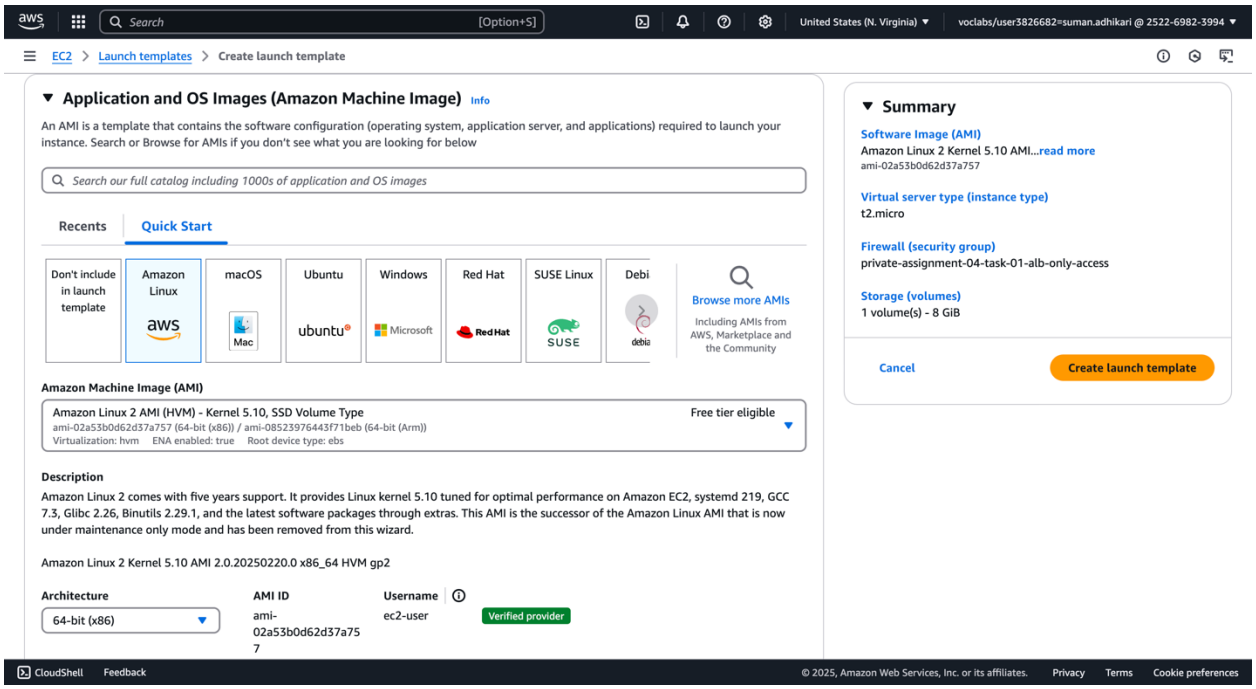

Create launch template

Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched at a later time. Templates can have multiple versions.

Launch template name and description

Launch template name - *required*

assignment-04-task-03-launch-template

Must be unique to this account. Max 128 chars. No spaces or special characters like '&', '*', '@'.

Template version description

A template for launching EC2 for assignment-04-task-03

Max 255 chars

Auto Scaling guidance | Info

Select this if you intend to use this template with EC2 Auto Scaling

☐ Provide guidance to help me set up a template that I can use with EC2 Auto Scaling

▸ **Template tags**

▸ **Source template**

Summary

Software Image (AMI)
Amazon Linux 2 Kernel 5.10 AMI...read more
ami-02a53b0d62d37a757

Virtual server type (instance type)
t2.micro

Firewall (security group)
private-assignment-04-task-01-alb-only-access

Storage (volumes)
1 volume(s) - 8 GiB

Cancel | Create launch template

Launch template contents

Specify the details of your launch template below. Leaving a field blank will result in the field not being included in the launch template.

Application and OS Images (Amazon Machine Image) Info

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Metadata response hop limit | Info

2

Allow tags in metadata | Info

Don't include in launch template

User data - *optional* | Info

Upload a file with your user data or enter it in the field.

Choose file

```
#!/bin/bash
yum install -y httpd
systemctl start httpd
systemctl enable httpd
echo "<h1>Hello Suman from $(hostname -f)</h1>" > /var/www/html/index.html
```

☐ User data has already been base64 encoded

Summary

Software Image (AMI)
Amazon Linux 2 Kernel 5.10 AMI...read more
ami-02a53b0d62d37a757

Virtual server type (instance type)
t2.micro

Firewall (security group)
private-assignment-04-task-01-alb-only-access

Storage (volumes)
1 volume(s) - 8 GiB

Cancel | Create template version

2. Create the Auto Scaling group (ASG) with above launch template.

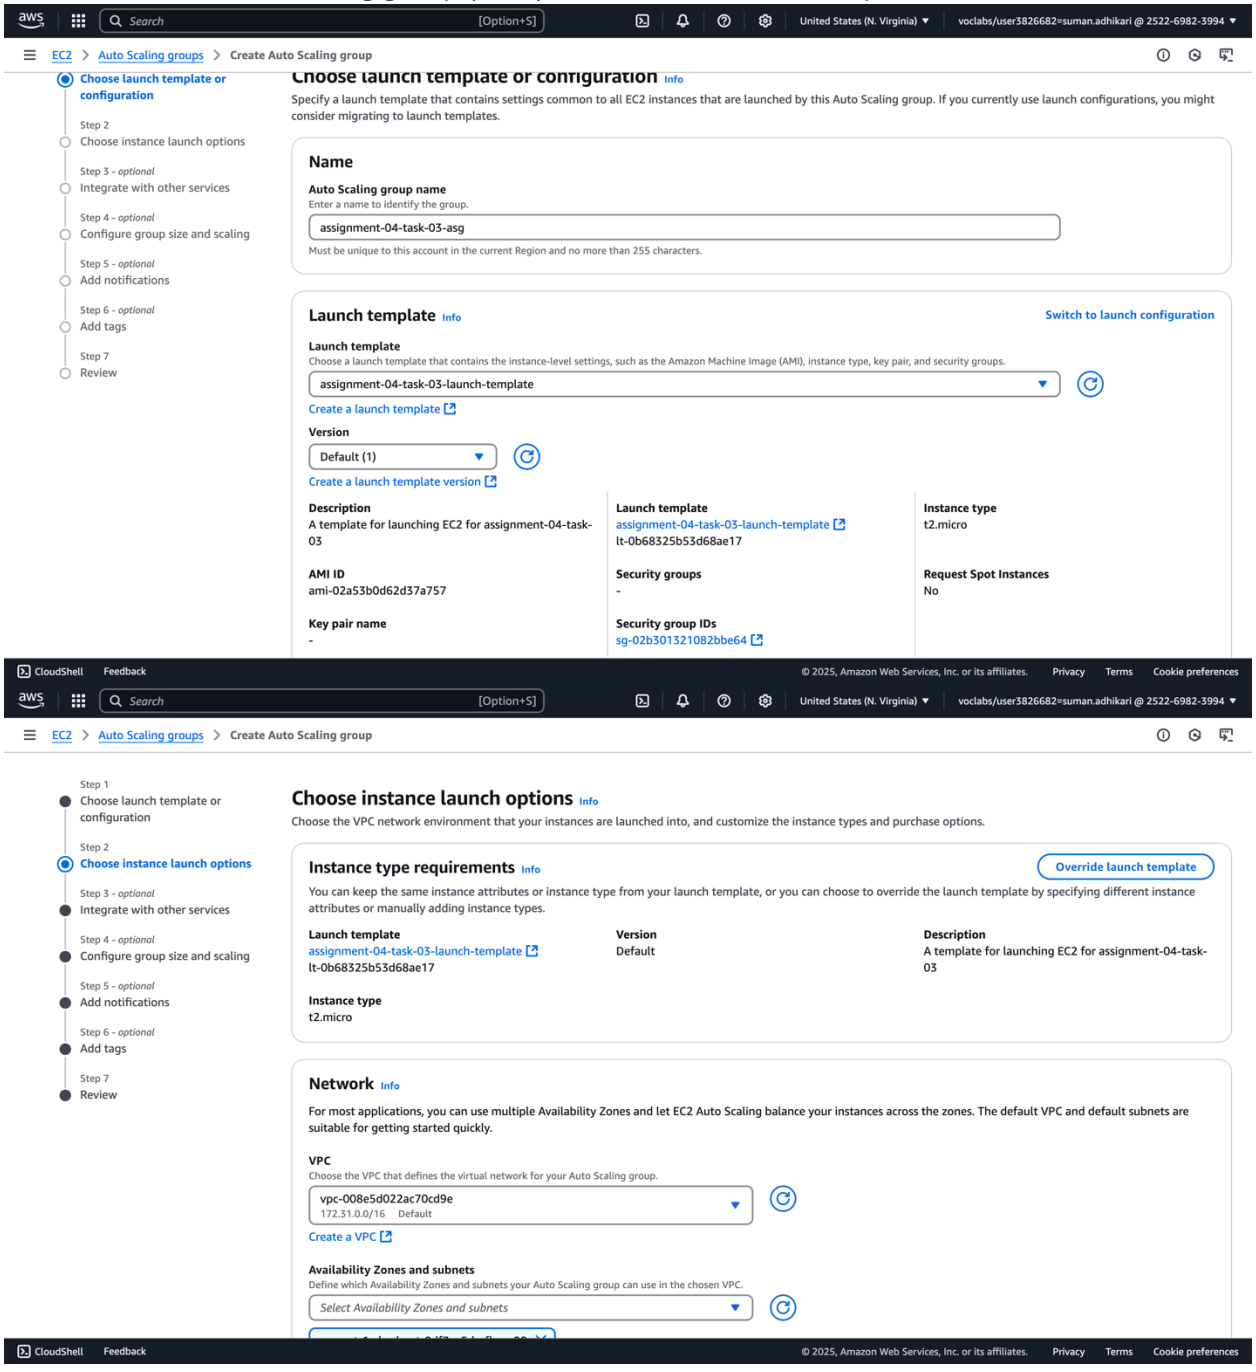

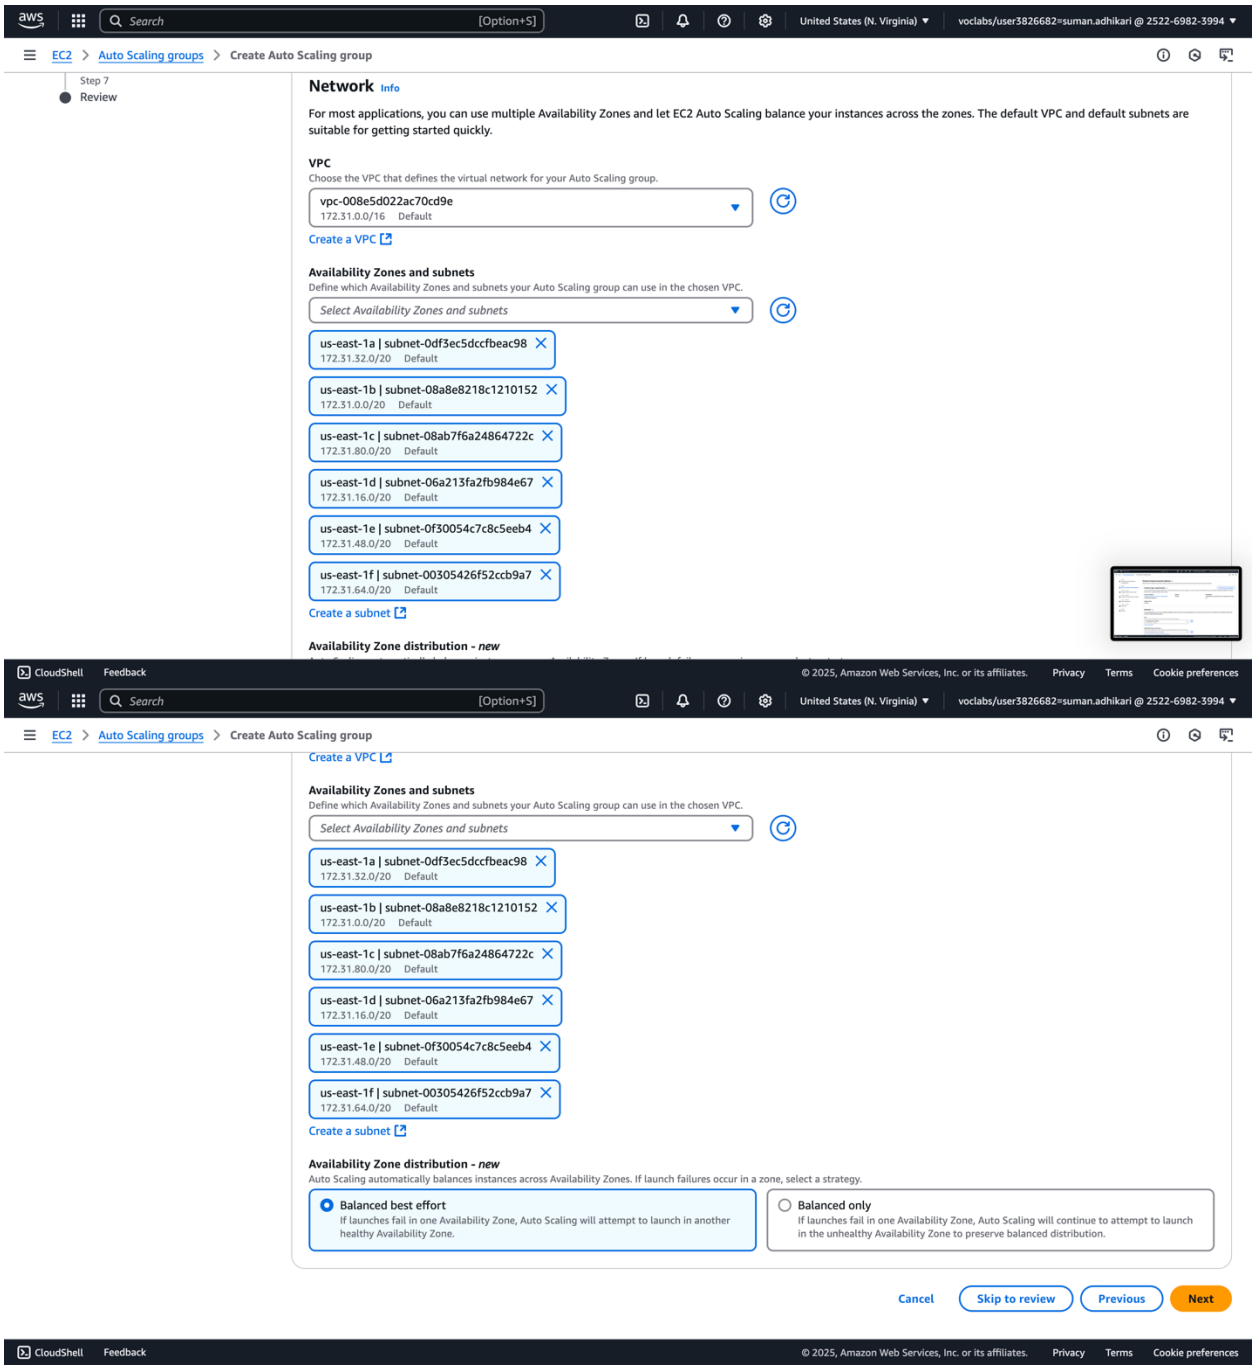
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 7
Review
Network Info
For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.
VPC
Choose the VPC that defines the virtual network for your Auto Scaling group.
vpc-008e5d022ac70cd9e
172.31.0.0/16 Default
Create a VPC
Availability Zones and subnets
Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.
Select Availability Zones and subnets
us-east-1a | subnet-0df3ec5dccfbeac98
172.31.32.0/20 Default
us-east-1b | subnet-08a8e8218c1210152
172.31.0.0/20 Default
us-east-1c | subnet-08ab7f6a24864722c
172.31.80.0/20 Default
us-east-1d | subnet-06a213fa2fb984e67
172.31.16.0/20 Default
us-east-1e | subnet-0f30054c7c8c5eeb4
172.31.48.0/20 Default
us-east-1f | subnet-00305426f52ccb9a7
172.31.64.0/20 Default
Create a subnet
Availability Zone distribution - new
CloudShell Feedback
© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > Create Auto Scaling group
Create a VPC
Availability Zones and subnets
Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.
Select Availability Zones and subnets
us-east-1a | subnet-0df3ec5dccfbeac98
172.31.32.0/20 Default
us-east-1b | subnet-08a8e8218c1210152
172.31.0.0/20 Default
us-east-1c | subnet-08ab7f6a24864722c
172.31.80.0/20 Default
us-east-1d | subnet-06a213fa2fb984e67
172.31.16.0/20 Default
us-east-1e | subnet-0f30054c7c8c5eeb4
172.31.48.0/20 Default
us-east-1f | subnet-00305426f52ccb9a7
172.31.64.0/20 Default
Create a subnet
Availability Zone distribution - new
Auto Scaling automatically balances instances across Availability Zones. If launch failures occur in a zone, select a strategy.
Balanced best effort
If launches fail in one Availability Zone, Auto Scaling will attempt to launch in another healthy Availability Zone.
Balanced only
If launches fail in one Availability Zone, Auto Scaling will continue to attempt to launch in the unhealthy Availability Zone to preserve balanced distribution.
Cancel
Skip to review
Previous
Next
CloudShell Feedback
© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

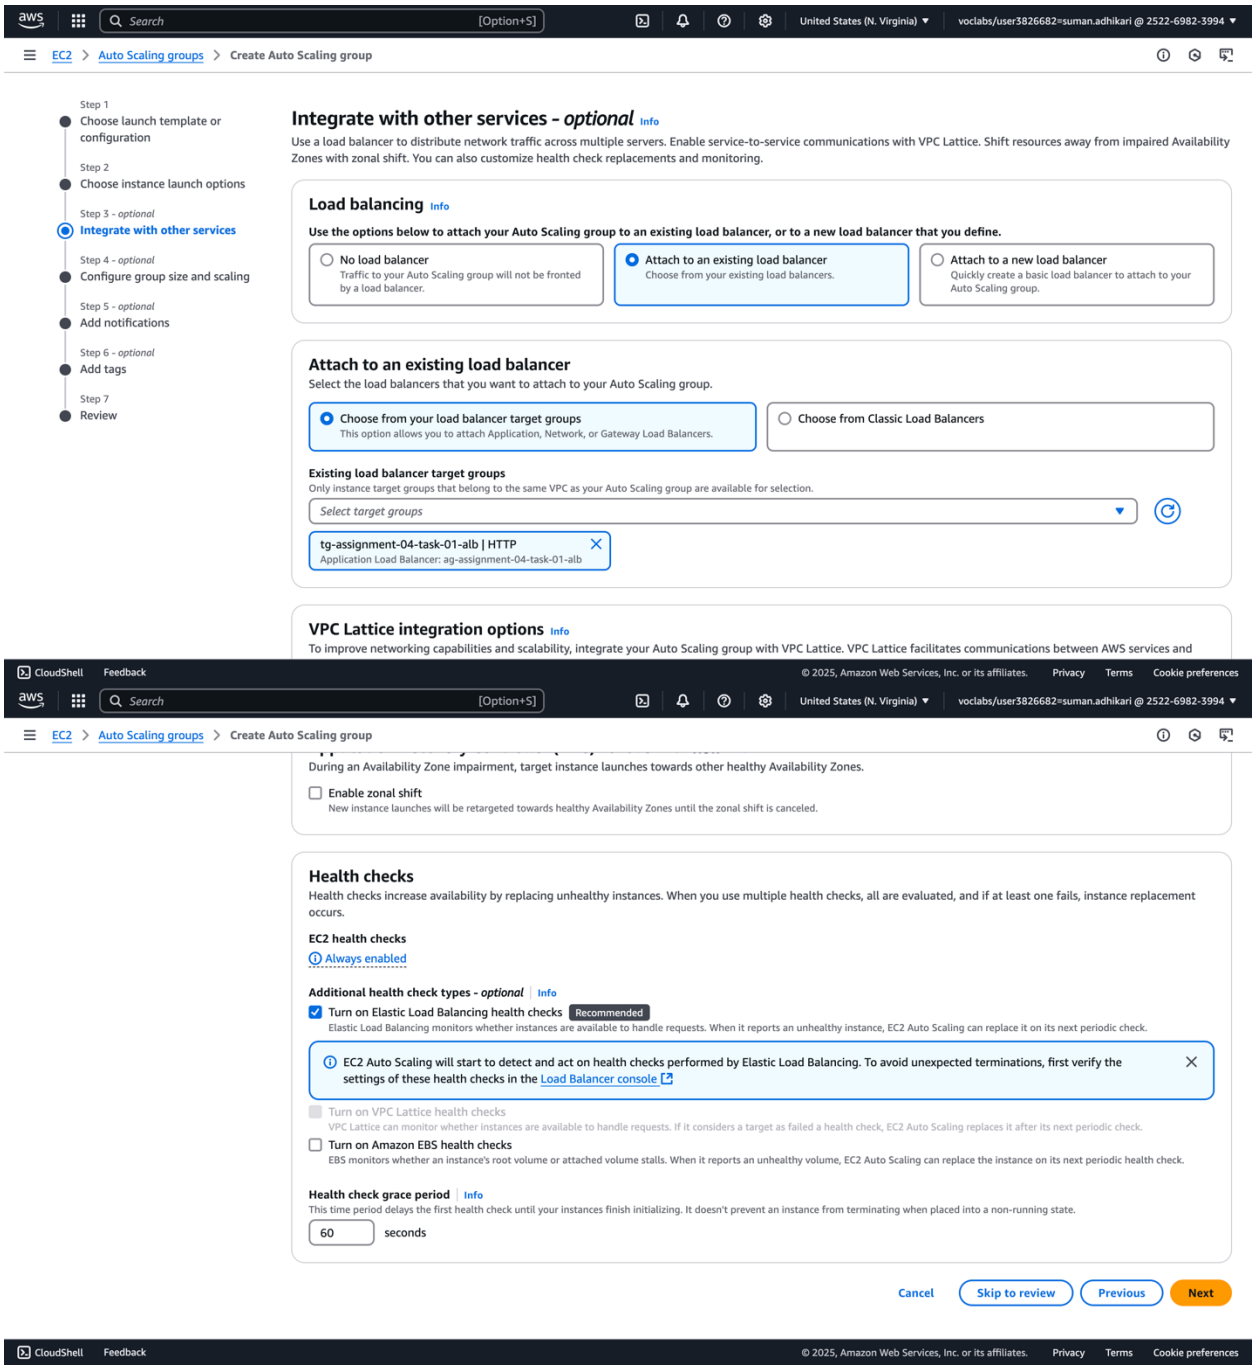
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1
Choose launch template or configuration
Step 2
Choose instance launch options
Step 3 - optional
Integrate with other services
Step 4 - optional
Configure group size and scaling
Step 5 - optional
Add notifications
Step 6 - optional
Add tags
Step 7
Review
Integrate with other services - optional Info
Use a load balancer to distribute network traffic across multiple servers. Enable service-to-service communications with VPC Lattice. Shift resources away from impaired Availability Zones with zonal shift. You can also customize health check replacements and monitoring.
Load balancing Info
Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.
No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.
Attach to an existing load balancer
Choose from your existing load balancers.
Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.
Attach to an existing load balancer
Select the load balancers that you want to attach to your Auto Scaling group.
Choose from your load balancer target groups
This option allows you to attach Application, Network, or Gateway Load Balancers.
Choose from Classic Load Balancers
Existing load balancer target groups
Only instance target groups that belong to the same VPC as your Auto Scaling group are available for selection.
Select target groups
tg-assignment-04-task-01-alb | HTTP
Application Load Balancer: ag-assignment-04-task-01-alb
VPC Lattice integration options Info
To improve networking capabilities and scalability, integrate your Auto Scaling group with VPC Lattice. VPC Lattice facilitates communications between AWS services and
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences
During an Availability Zone impairment, target instance launches towards other healthy Availability Zones.
Enable zonal shift
New instance launches will be retargeted towards healthy Availability Zones until the zonal shift is canceled.
Health checks
Health checks increase availability by replacing unhealthy instances. When you use multiple health checks, all are evaluated, and if at least one fails, instance replacement occurs.
EC2 health checks
Always enabled
Additional health check types - optional Info
Turn on Elastic Load Balancing health checks
Recommended
Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.
EC2 Auto Scaling will start to detect and act on health checks performed by Elastic Load Balancing. To avoid unexpected terminations, first verify the settings of these health checks in the Load Balancer console
Turn on VPC Lattice health checks
VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.
Turn on Amazon EBS health checks
EBS monitors whether an instance's root volume or attached volume stalls. When it reports an unhealthy volume, EC2 Auto Scaling can replace the instance on its next periodic health check.
Health check grace period Info
This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.
60
seconds
Cancel
Skip to review
Previous
Next

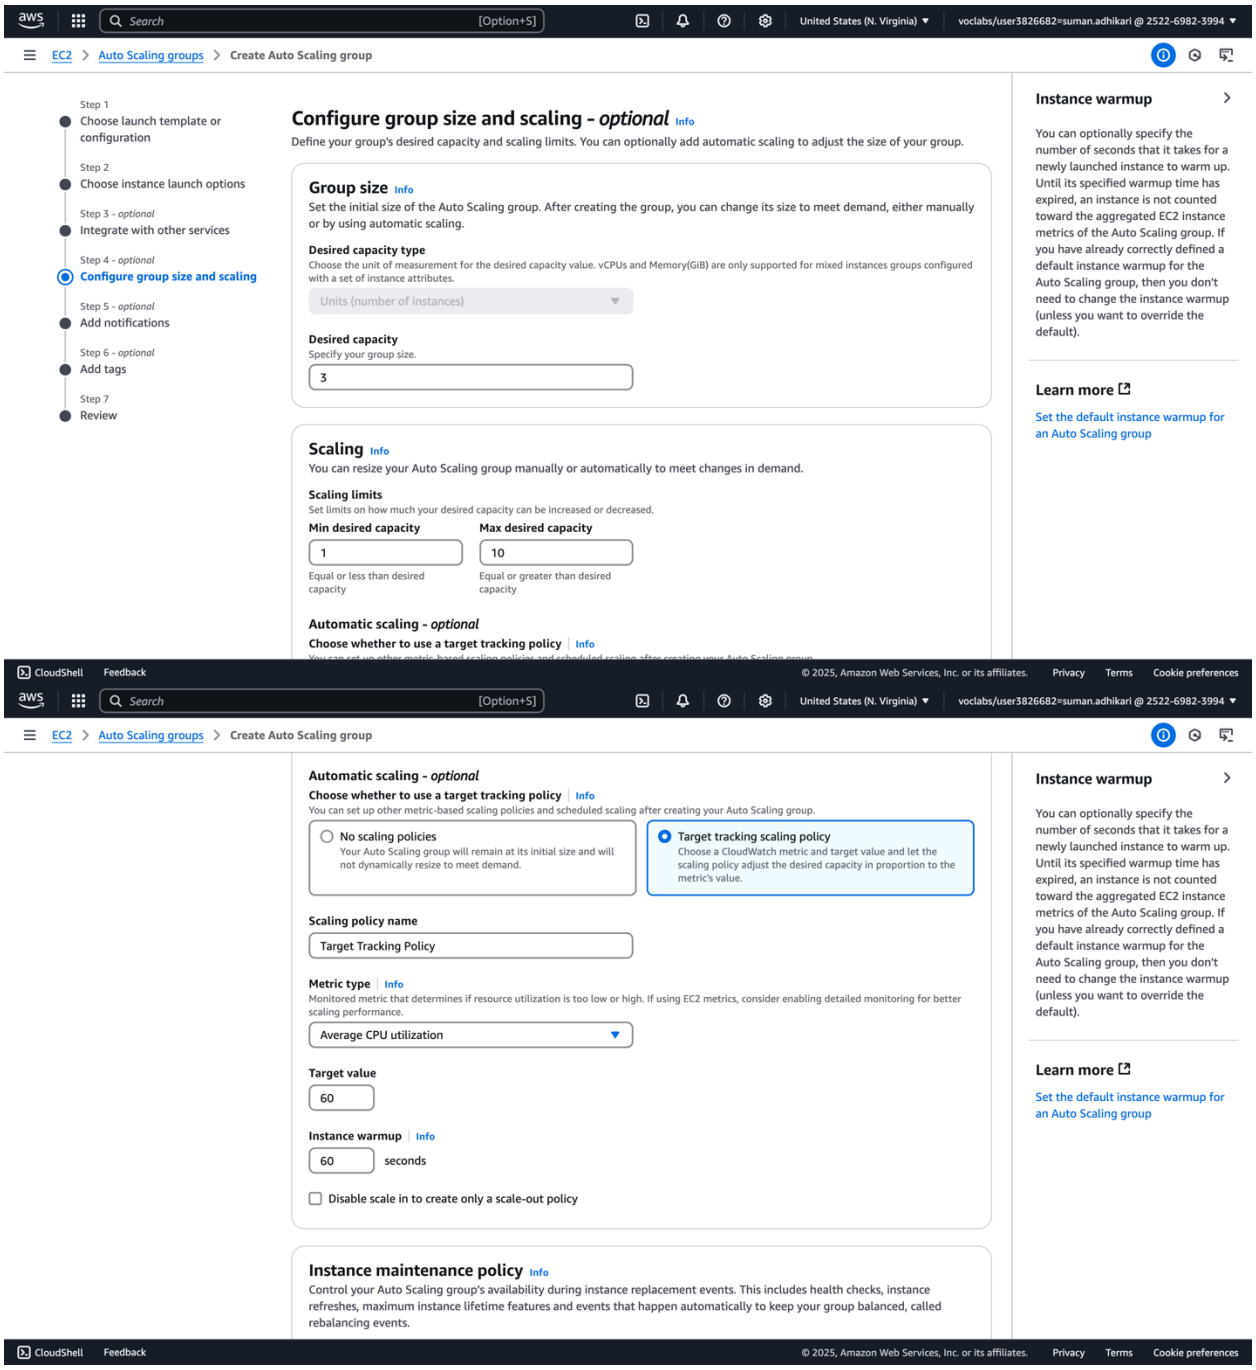
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1
Choose launch template or configuration
Step 2
Choose instance launch options
Step 3 - optional
Integrate with other services
Step 4 - optional
Configure group size and scaling
Step 5 - optional
Add notifications
Step 6 - optional
Add tags
Step 7
Review
Configure group size and scaling - optional Info
Define your group's desired capacity and scaling limits. You can optionally add automatic scaling to adjust the size of your group.
Group size Info
Set the initial size of the Auto Scaling group. After creating the group, you can change its size to meet demand, either manually or by using automatic scaling.
Desired capacity type
Choose the unit of measurement for the desired capacity value. vCPUs and Memory(GiB) are only supported for mixed instances groups configured with a set of instance attributes.
Units (number of instances)
Desired capacity
Specify your group size.
3
Scaling Info
You can resize your Auto Scaling group manually or automatically to meet changes in demand.
Scaling limits
Set limits on how much your desired capacity can be increased or decreased.
Min desired capacity
1
Equal or less than desired capacity
Max desired capacity
10
Equal or greater than desired capacity
Automatic scaling - optional
Choose whether to use a target tracking policy Info
Instance warmup
You can optionally specify the number of seconds that it takes for a newly launched instance to warm up. Until its specified warmup time has expired, an instance is not counted toward the aggregated EC2 instance metrics of the Auto Scaling group. If you have already correctly defined a default instance warmup for the Auto Scaling group, then you don't need to change the instance warmup (unless you want to override the default).
Learn more
Set the default instance warmup for an Auto Scaling group
CloudShell Feedback
© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > Create Auto Scaling group
Automatic scaling - optional
Choose whether to use a target tracking policy Info
You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.
No scaling policies
Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.
Target tracking scaling policy
Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.
Scaling policy name
Target Tracking Policy
Metric type Info
Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.
Average CPU utilization
Target value
60
Instance warmup Info
60
seconds
Disable scale in to create only a scale-out policy
Instance maintenance policy Info
Control your Auto Scaling group's availability during instance replacement events. This includes health checks, instance refreshes, maximum instance lifetime features and events that happen automatically to keep your group balanced, called rebalancing events.
Instance warmup
You can optionally specify the number of seconds that it takes for a newly launched instance to warm up. Until its specified warmup time has expired, an instance is not counted toward the aggregated EC2 instance metrics of the Auto Scaling group. If you have already correctly defined a default instance warmup for the Auto Scaling group, then you don't need to change the instance warmup (unless you want to override the default).
Learn more
Set the default instance warmup for an Auto Scaling group
CloudShell Feedback
© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

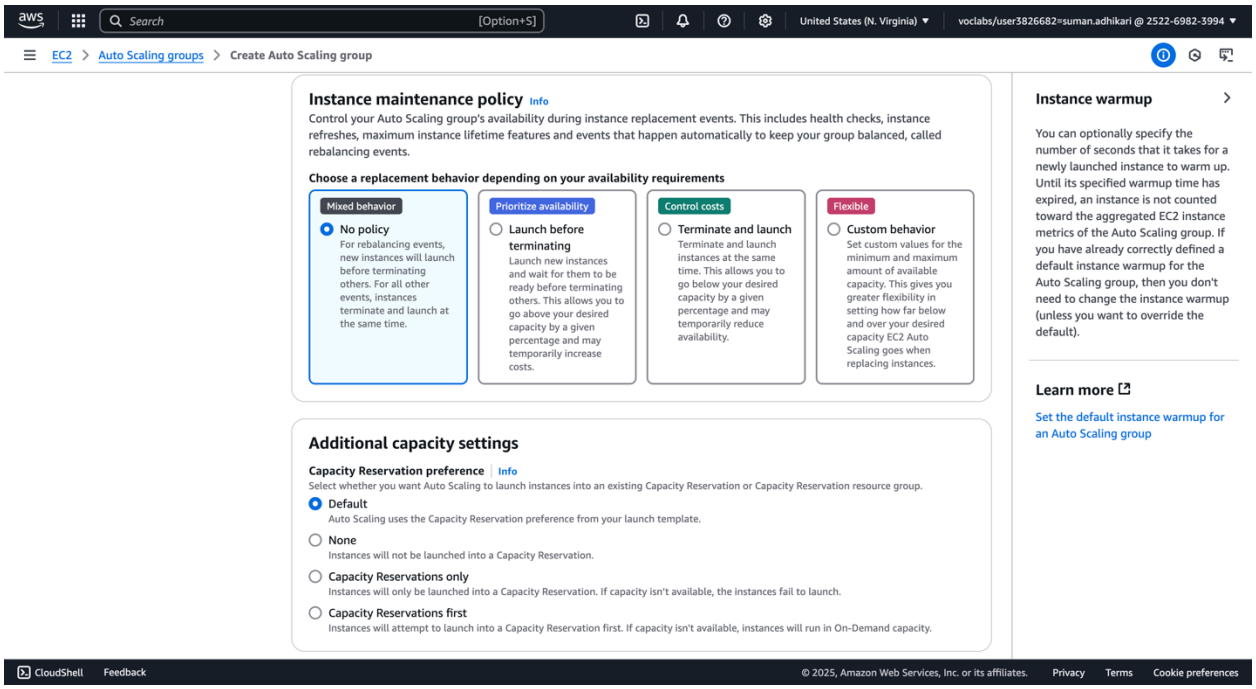

3. Review the ASG settings.

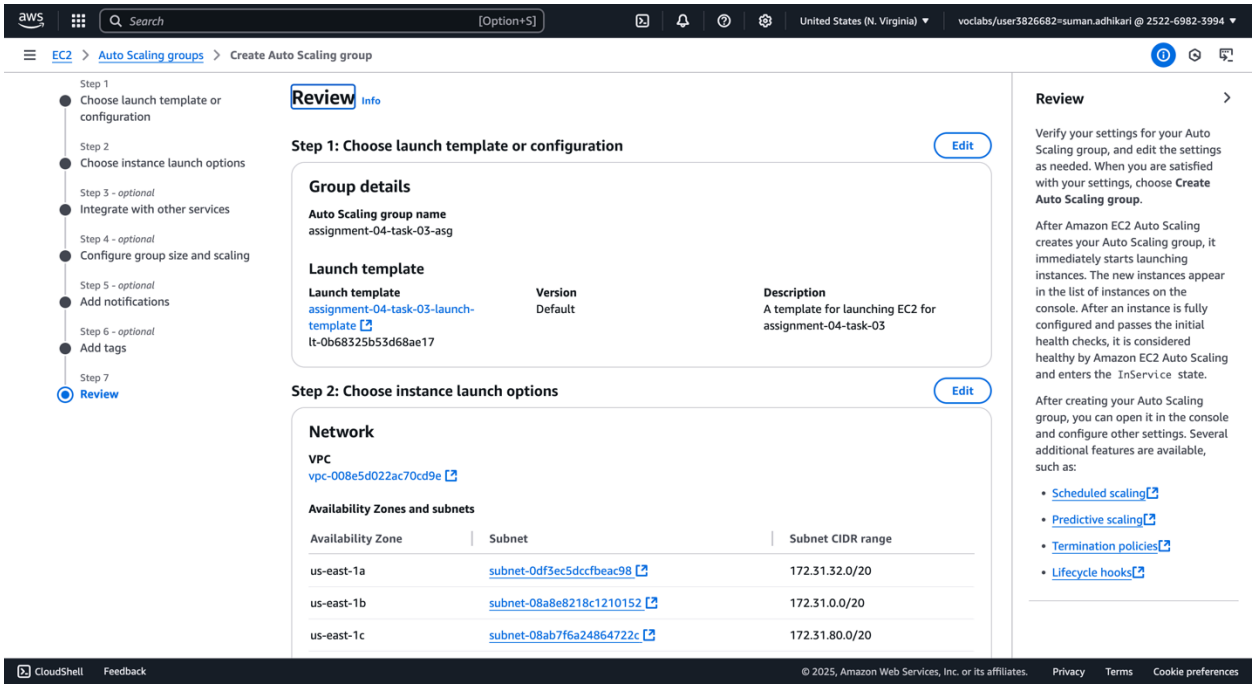

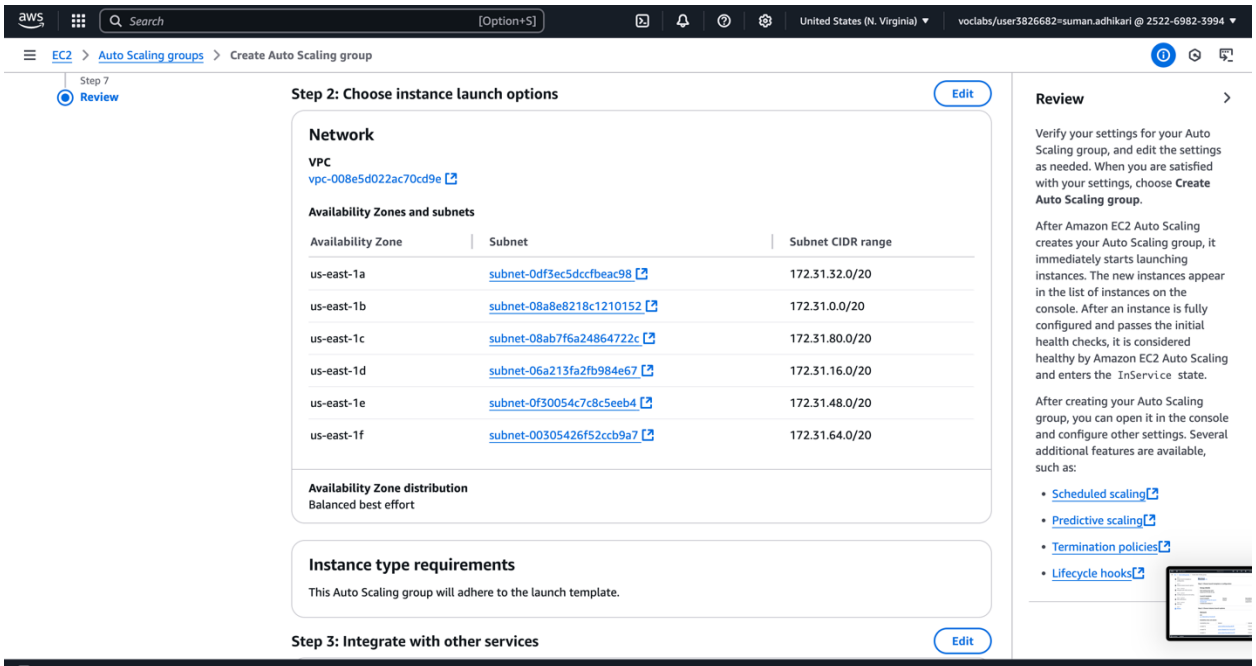
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 7
Review
Step 2: Choose instance launch options
Edit
Network
VPC
vpc-008e5d022ac70cd9e
Availability Zones and subnets
Availability Zone | Subnet | Subnet CIDR range
us-east-1a | subnet-0df3ec5dccfbeac98 | 172.31.32.0/20
us-east-1b | subnet-08a8e8218c1210152 | 172.31.0.0/20
us-east-1c | subnet-08ab7f6a24864722c | 172.31.80.0/20
us-east-1d | subnet-06a213fa2fb984e67 | 172.31.16.0/20
us-east-1e | subnet-0f30054c7c8c5eeb4 | 172.31.48.0/20
us-east-1f | subnet-00305426f52ccb9a7 | 172.31.64.0/20
Availability Zone distribution
Balanced best effort
Instance type requirements
This Auto Scaling group will adhere to the launch template.
Step 3: Integrate with other services
Edit
Review
Verify your settings for your Auto Scaling group, and edit the settings as needed. When you are satisfied with your settings, choose Create Auto Scaling group.
After Amazon EC2 Auto Scaling creates your Auto Scaling group, it immediately starts launching instances. The new instances appear in the list of instances on the console. After an instance is fully configured and passes the initial health checks, it is considered healthy by Amazon EC2 Auto Scaling and enters the InService state.
After creating your Auto Scaling group, you can open it in the console and configure other settings. Several additional features are available, such as:
Scheduled scaling
Predictive scaling
Termination policies
Lifecycle hooks
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

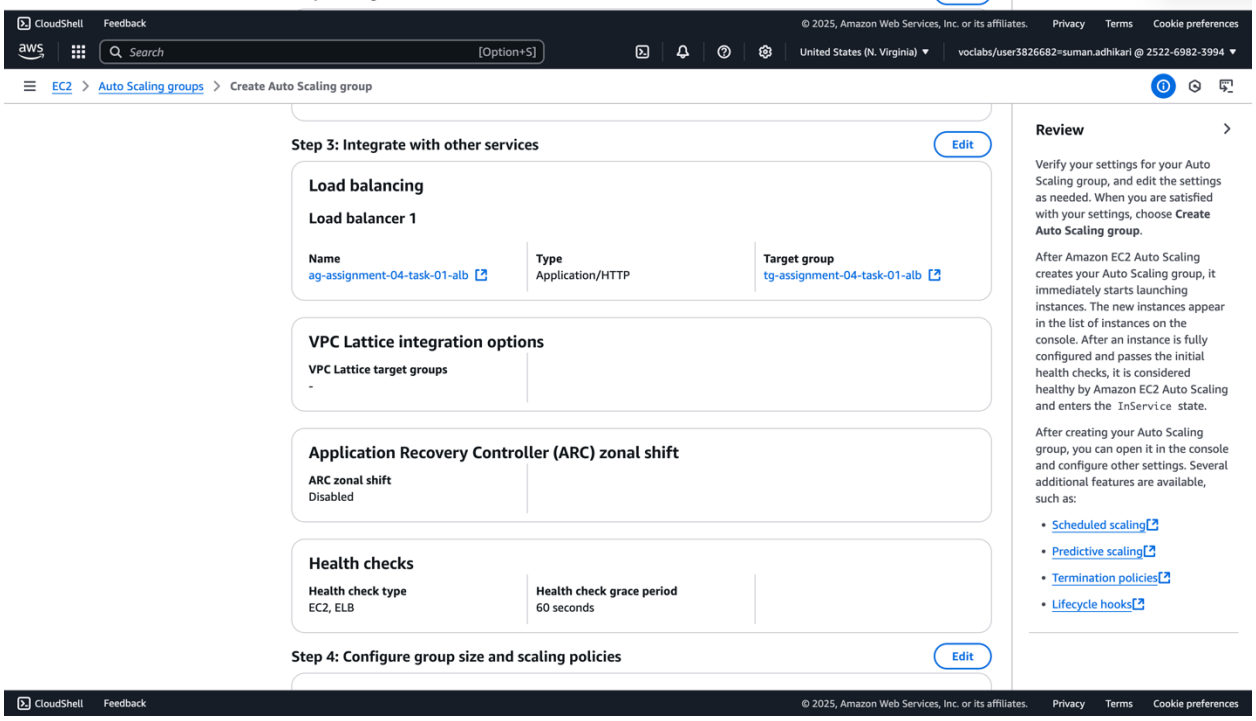
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 3: Integrate with other services
Edit
Load balancing
Load balancer 1
Name
ag-assignment-04-task-01-alb
Type
Application/HTTP
Target group
tg-assignment-04-task-01-alb
VPC Lattice integration options
VPC Lattice target groups
-
Application Recovery Controller (ARC) zonal shift
ARC zonal shift
Disabled
Health checks
Health check type
EC2, ELB
Health check grace period
60 seconds
Step 4: Configure group size and scaling policies
Edit
Review
Verify your settings for your Auto Scaling group, and edit the settings as needed. When you are satisfied with your settings, choose Create Auto Scaling group.
After Amazon EC2 Auto Scaling creates your Auto Scaling group, it immediately starts launching instances. The new instances appear in the list of instances on the console. After an instance is fully configured and passes the initial health checks, it is considered healthy by Amazon EC2 Auto Scaling and enters the InService state.
After creating your Auto Scaling group, you can open it in the console and configure other settings. Several additional features are available, such as:
Scheduled scaling
Predictive scaling
Termination policies
Lifecycle hooks
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

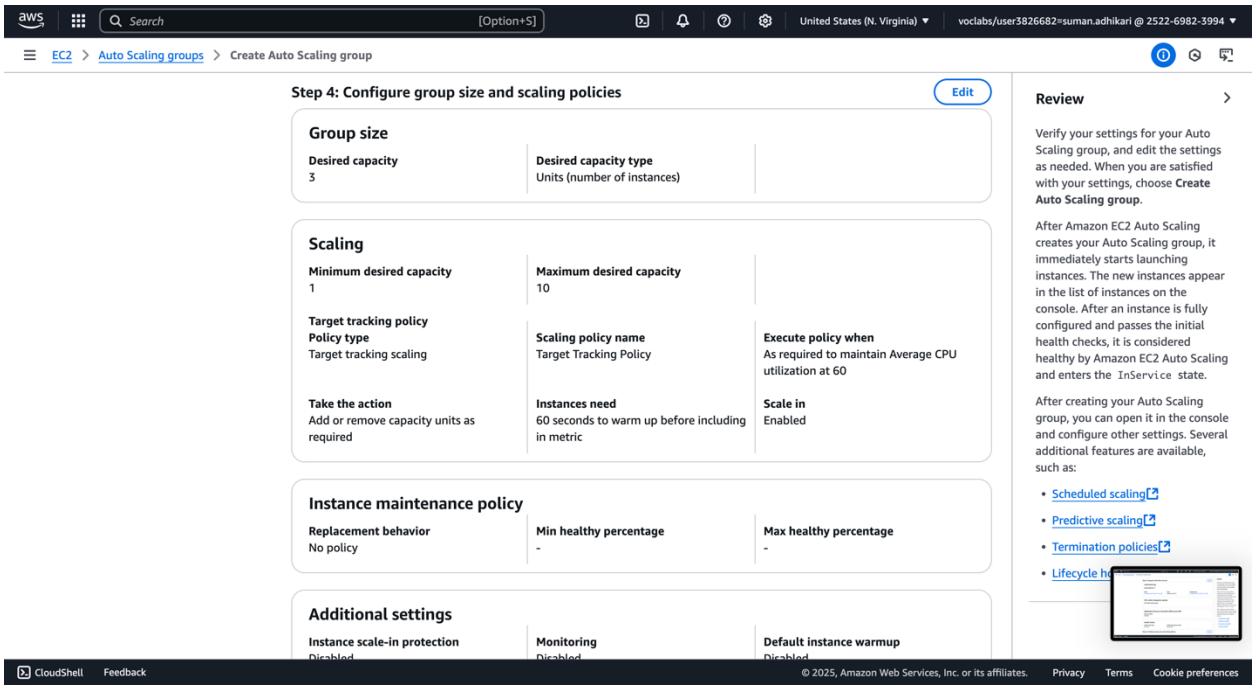

4. ASG is now active and here is the initial activity.

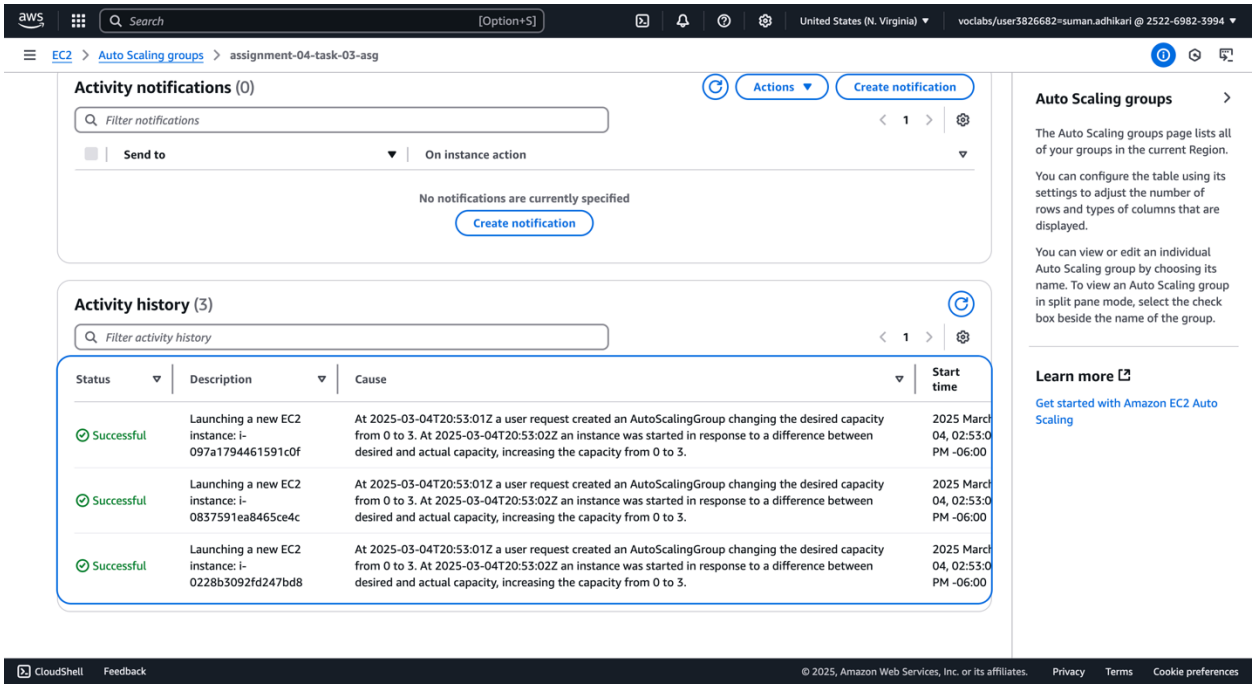

5. Add SSH access to instance for the stress test now.

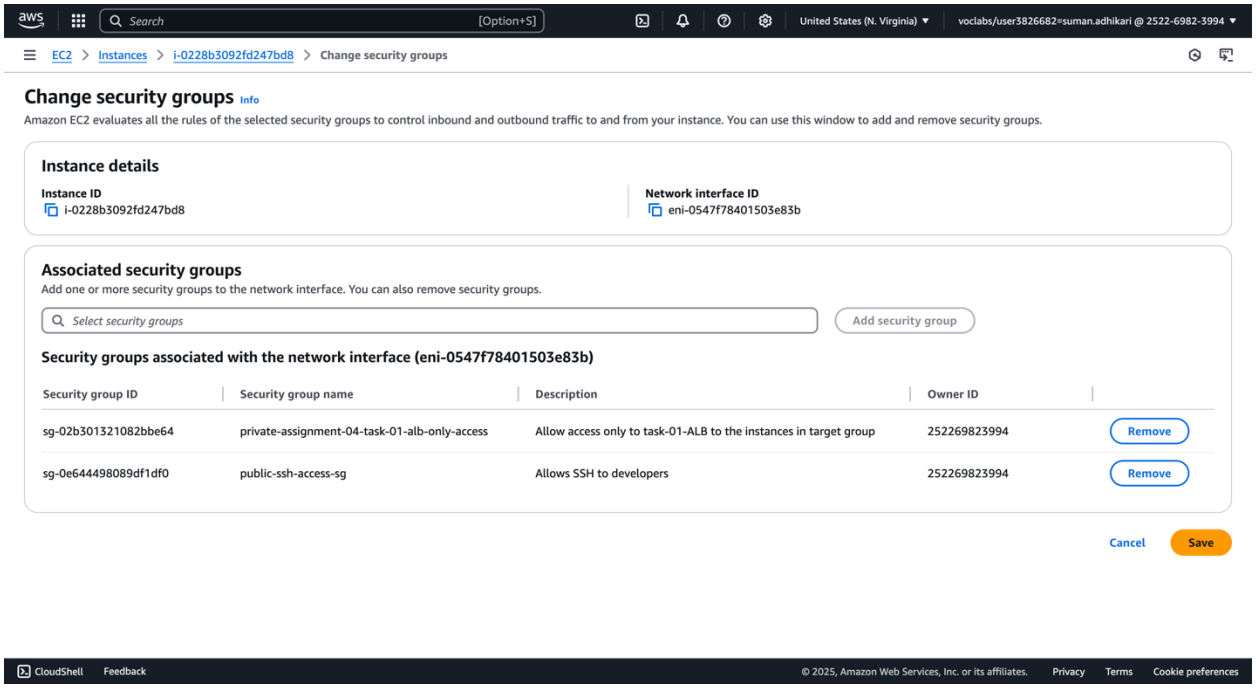

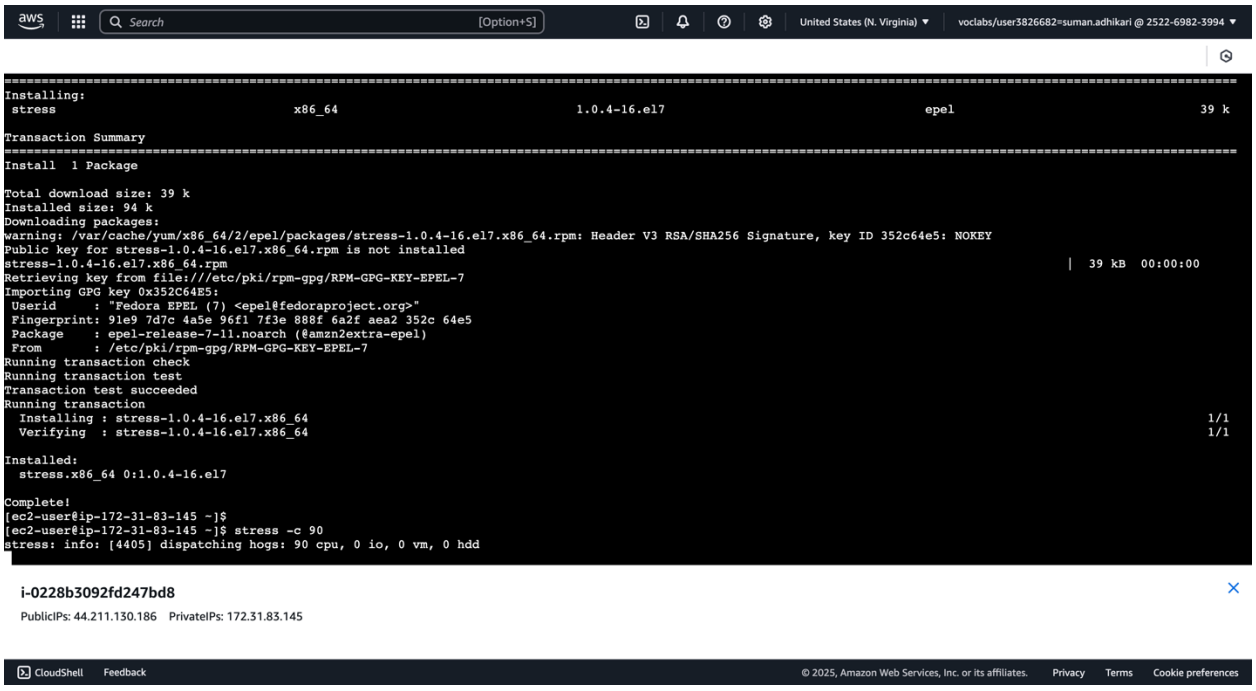

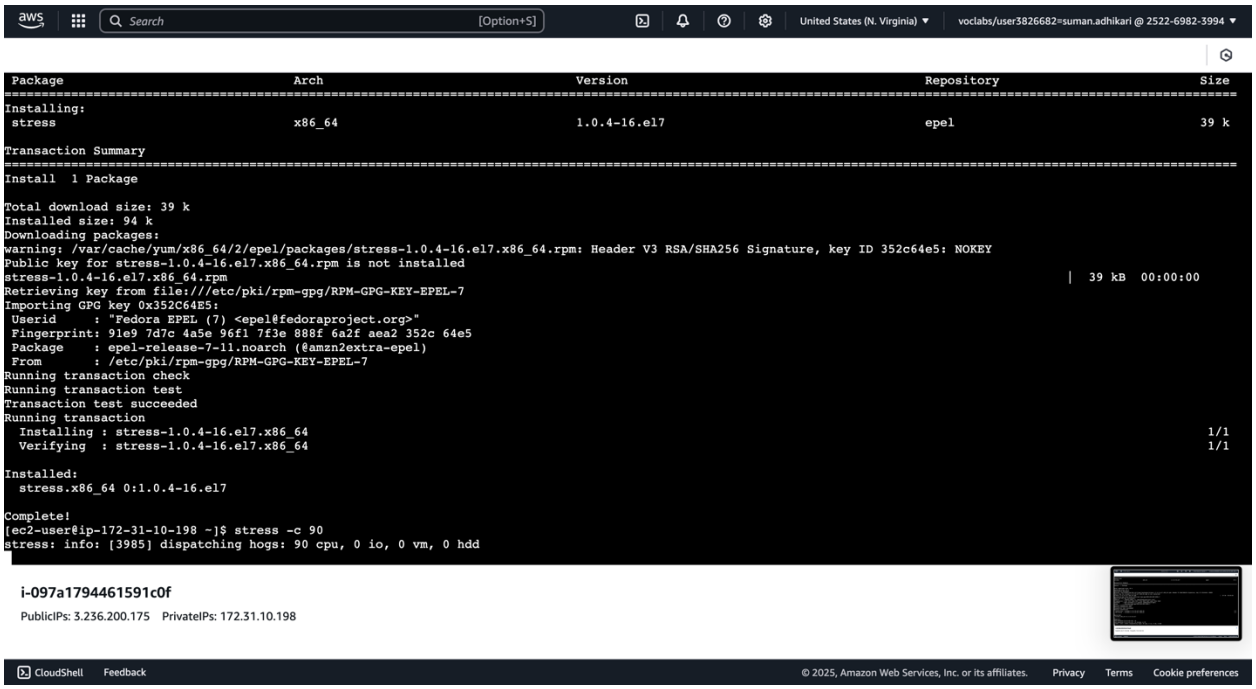

6. Check activity now for updates to ASG.

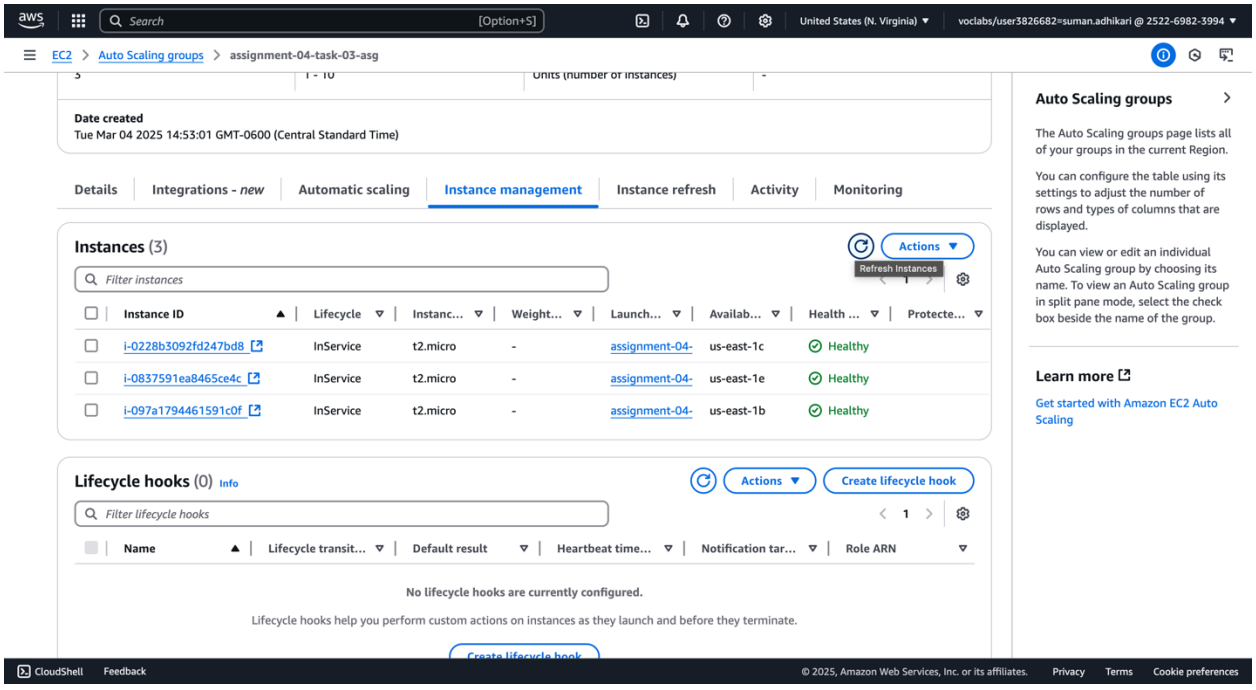

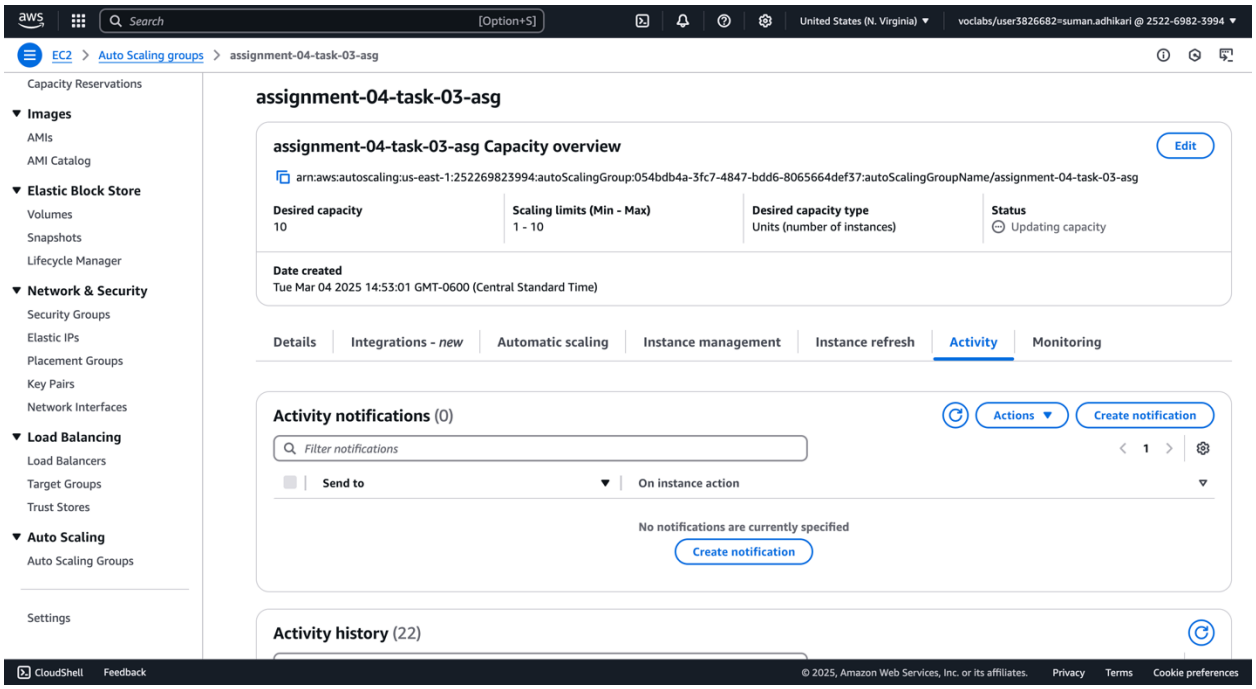
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > assignment-04-task-03-asg
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
Load Balancing
Load Balancers
Target Groups
Trust Stores
Auto Scaling
Auto Scaling Groups
Settings
assignment-04-task-03-asg
assignment-04-task-03-asg Capacity overview
Edit
arn:aws:autoscaling:us-east-1:252269823994:autoScalingGroup:054bdb4a-3fc7-4847-bdd6-8065664def37:autoScalingGroupName/assignment-04-task-03-asg
Desired capacity
10
Scaling limits (Min - Max)
1 - 10
Desired capacity type
Units (number of instances)
Status
Updating capacity
Date created
Tue Mar 04 2025 14:53:01 GMT-0600 (Central Standard Time)
Details
Integrations - new
Automatic scaling
Instance management
Instance refresh
Activity
Monitoring
Activity notifications (0)
Actions
Create notification
Filter notifications
1
Send to
On instance action
No notifications are currently specified
Create notification
Activity history (22)
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

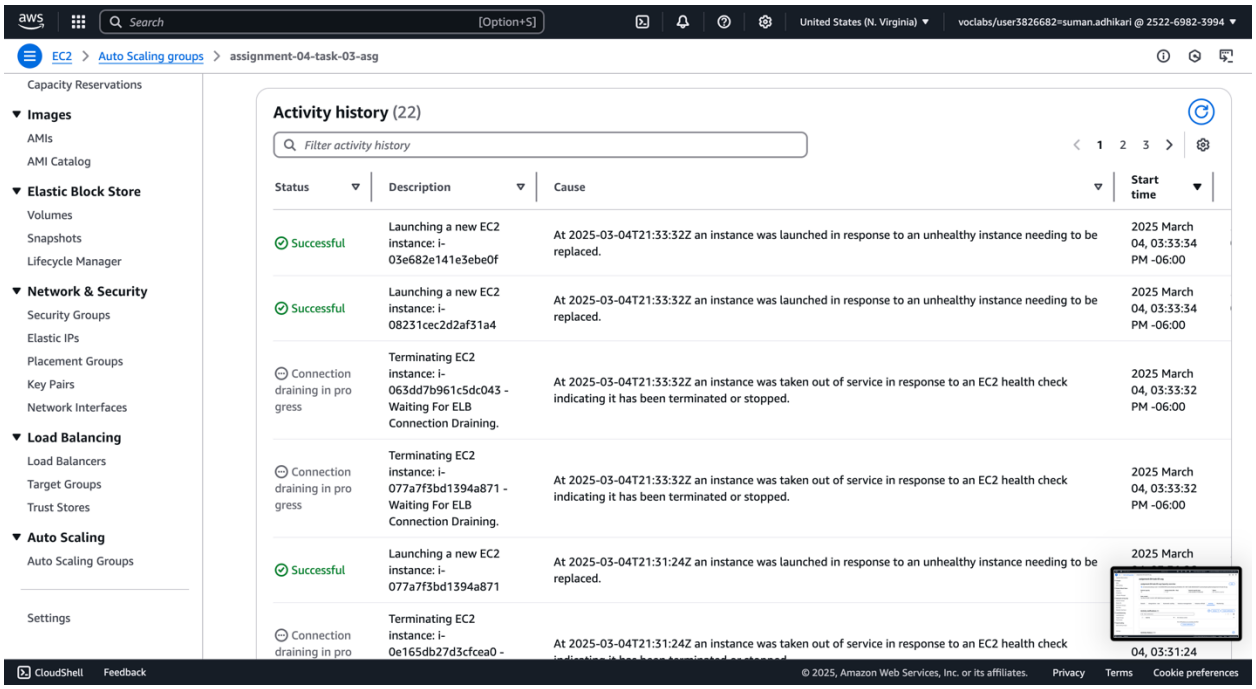
Search [Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > assignment-04-task-03-asg
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
Load Balancing
Load Balancers
Target Groups
Trust Stores
Auto Scaling
Auto Scaling Groups
Settings
Activity history (22)
Filter activity history
1 2 3
Status
Description
Cause
Start time
Successful
Launching a new EC2 instance: i-03e682e141e3ebe0f
At 2025-03-04T21:33:32Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March 04, 03:33:34 PM -06:00
Successful
Launching a new EC2 instance: i-08231cec2d2af31a4
At 2025-03-04T21:33:32Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March 04, 03:33:34 PM -06:00
Connection draining in progress
Terminating EC2 instance: i-063dd7b961c5dc043 - Waiting For ELB Connection Draining.
At 2025-03-04T21:33:32Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.
2025 March 04, 03:33:32 PM -06:00
Connection draining in progress
Terminating EC2 instance: i-077a7f3bd1394a871 - Waiting For ELB Connection Draining.
At 2025-03-04T21:33:32Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.
2025 March 04, 03:33:32 PM -06:00
Successful
Launching a new EC2 instance: i-077a7f3bd1394a871
At 2025-03-04T21:31:24Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March
Connection draining in pro
Terminating EC2 instance: i-0e165db27d3cfcea0 -
At 2025-03-04T21:31:24Z an instance was taken out of service in response to an EC2 health check
04, 03:31:24
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

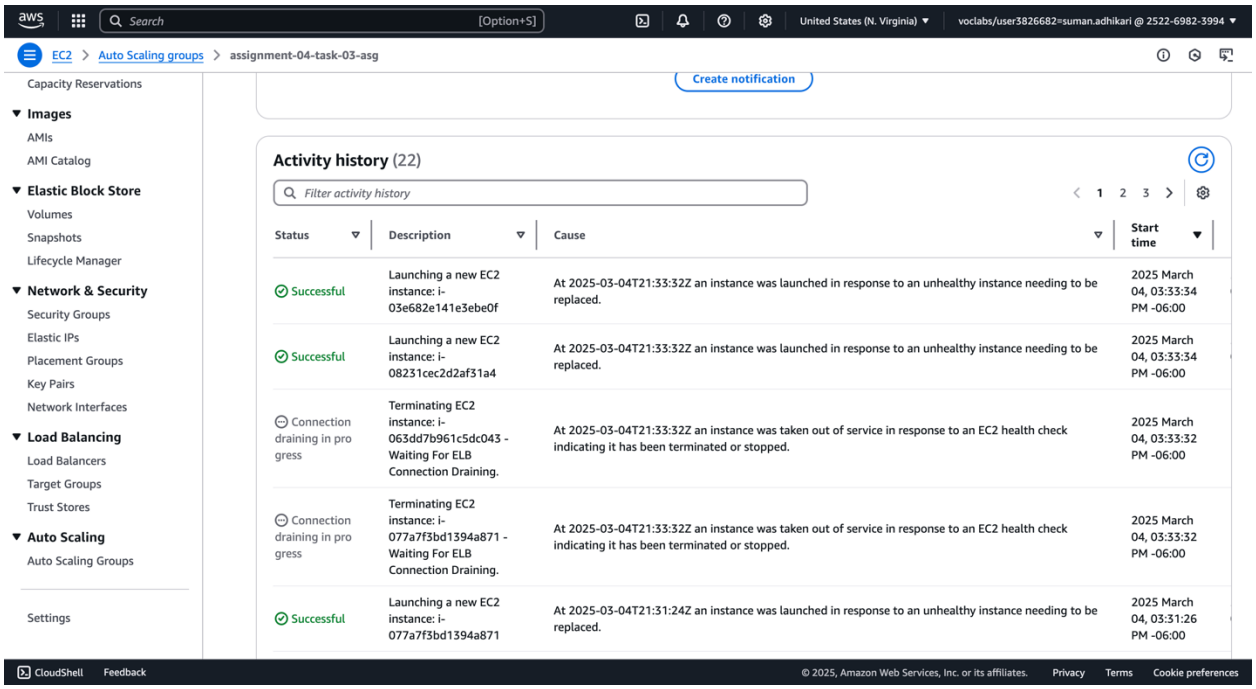
aws
Search
[Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > assignment-04-task-03-asg
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
Load Balancing
Load Balancers
Target Groups
Trust Stores
Auto Scaling
Auto Scaling Groups
Settings
Create notification
Activity history (22)
Filter activity history
1 2 3
Status
Description
Cause
Start time
Successful
Launching a new EC2 instance: i-03e682e141e3ebe0f
At 2025-03-04T21:33:32Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March 04, 03:33:34 PM -06:00
Successful
Launching a new EC2 instance: i-08231cec2d2af31a4
At 2025-03-04T21:33:32Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March 04, 03:33:34 PM -06:00
Connection draining in pro gress
Terminating EC2 instance: i-063dd7b961c5dc043 - Waiting For ELB Connection Draining.
At 2025-03-04T21:33:32Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.
2025 March 04, 03:33:32 PM -06:00
Connection draining in pro gress
Terminating EC2 instance: i-077a7f3bd1394a871 - Waiting For ELB Connection Draining.
At 2025-03-04T21:33:32Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.
2025 March 04, 03:33:32 PM -06:00
Successful
Launching a new EC2 instance: i-077a7f3bd1394a871
At 2025-03-04T21:31:24Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March 04, 03:31:26 PM -06:00
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

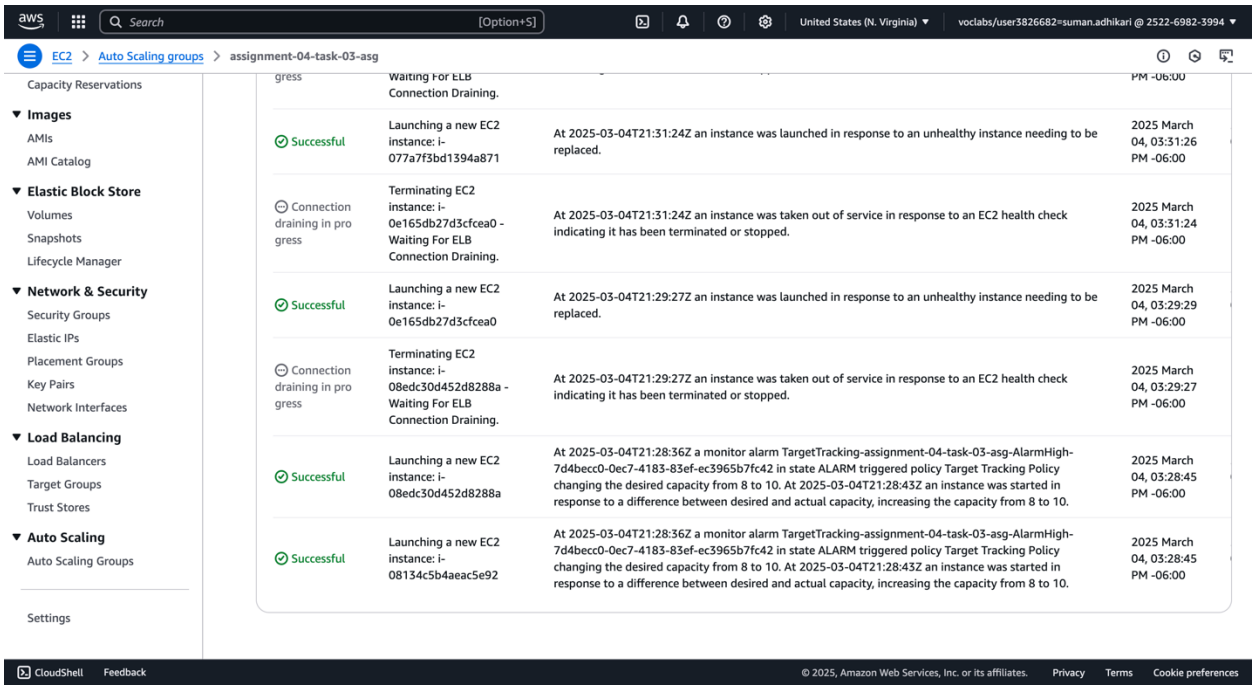
aws
Search
[Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > assignment-04-task-03-asg
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
Load Balancing
Load Balancers
Target Groups
Trust Stores
Auto Scaling
Auto Scaling Groups
Settings
Waiting For ELB Connection Draining.
Successful
Launching a new EC2 instance: i-077a7f3bd1394a871
At 2025-03-04T21:31:24Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March 04, 03:31:26 PM -06:00
Connection draining in pro gress
Terminating EC2 instance: i-0e165db27d3cfcea0 - Waiting For ELB Connection Draining.
At 2025-03-04T21:31:24Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.
2025 March 04, 03:31:24 PM -06:00
Successful
Launching a new EC2 instance: i-0e165db27d3cfcea0
At 2025-03-04T21:29:27Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March 04, 03:29:29 PM -06:00
Connection draining in pro gress
Terminating EC2 instance: i-08edc30d452d8288a - Waiting For ELB Connection Draining.
At 2025-03-04T21:29:27Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.
2025 March 04, 03:29:27 PM -06:00
Successful
Launching a new EC2 instance: i-08edc30d452d8288a
At 2025-03-04T21:28:36Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmHigh-7d4becc0-0ec7-4183-83ef-ec3965b7fc42 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 8 to 10. At 2025-03-04T21:28:43Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 8 to 10.
2025 March 04, 03:28:45 PM -06:00
Successful
Launching a new EC2 instance: i-08134c5b4aeac5e92
At 2025-03-04T21:28:36Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmHigh-7d4becc0-0ec7-4183-83ef-ec3965b7fc42 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 8 to 10. At 2025-03-04T21:28:43Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 8 to 10.
2025 March 04, 03:28:45 PM -06:00
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

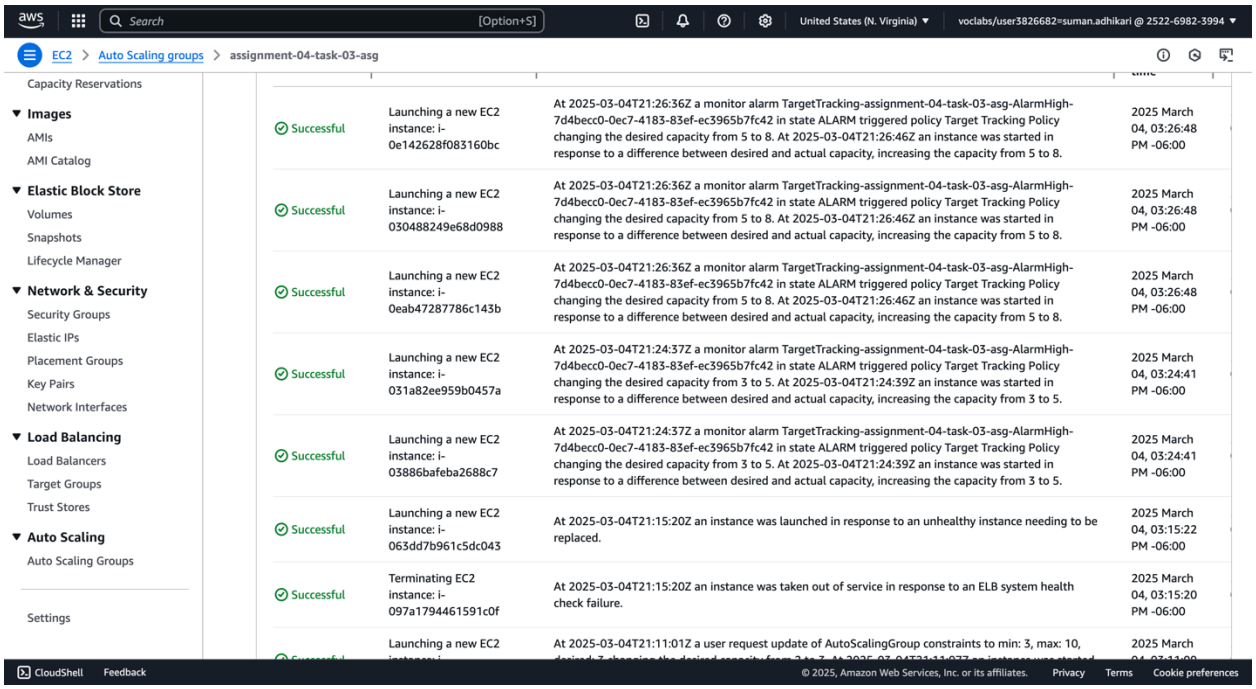
aws
Search
[Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > assignment-04-task-03-asg
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
Load Balancing
Load Balancers
Target Groups
Trust Stores
Auto Scaling
Auto Scaling Groups
Settings
Successful
Launching a new EC2 instance: i-0e142628f083160bc
At 2025-03-04T21:26:36Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmHigh-7d4becc0-0ec7-4183-83ef-ec3965b7fc42 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 5 to 8. At 2025-03-04T21:26:46Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 5 to 8.
2025 March 04, 03:26:48 PM -06:00
Successful
Launching a new EC2 instance: i-030488249e68d0988
At 2025-03-04T21:26:36Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmHigh-7d4becc0-0ec7-4183-83ef-ec3965b7fc42 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 5 to 8. At 2025-03-04T21:26:46Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 5 to 8.
2025 March 04, 03:26:48 PM -06:00
Successful
Launching a new EC2 instance: i-0eab47287786c143b
At 2025-03-04T21:26:36Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmHigh-7d4becc0-0ec7-4183-83ef-ec3965b7fc42 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 5 to 8. At 2025-03-04T21:26:46Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 5 to 8.
2025 March 04, 03:26:48 PM -06:00
Successful
Launching a new EC2 instance: i-031a82ee959b0457a
At 2025-03-04T21:24:37Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmHigh-7d4becc0-0ec7-4183-83ef-ec3965b7fc42 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 3 to 5. At 2025-03-04T21:24:39Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 3 to 5.
2025 March 04, 03:24:41 PM -06:00
Successful
Launching a new EC2 instance: i-03886bafeba2688c7
At 2025-03-04T21:24:37Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmHigh-7d4becc0-0ec7-4183-83ef-ec3965b7fc42 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 3 to 5. At 2025-03-04T21:24:39Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 3 to 5.
2025 March 04, 03:24:41 PM -06:00
Successful
Launching a new EC2 instance: i-063dd7b961c5dc043
At 2025-03-04T21:15:20Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March 04, 03:15:22 PM -06:00
Successful
Terminating EC2 instance: i-097a1794461591c0f
At 2025-03-04T21:15:20Z an instance was taken out of service in response to an ELB system health check failure.
2025 March 04, 03:15:20 PM -06:00
Successful
Launching a new EC2
At 2025-03-04T21:11:01Z a user request update of AutoScalingGroup constraints to min: 3, max: 10,
2025 March
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

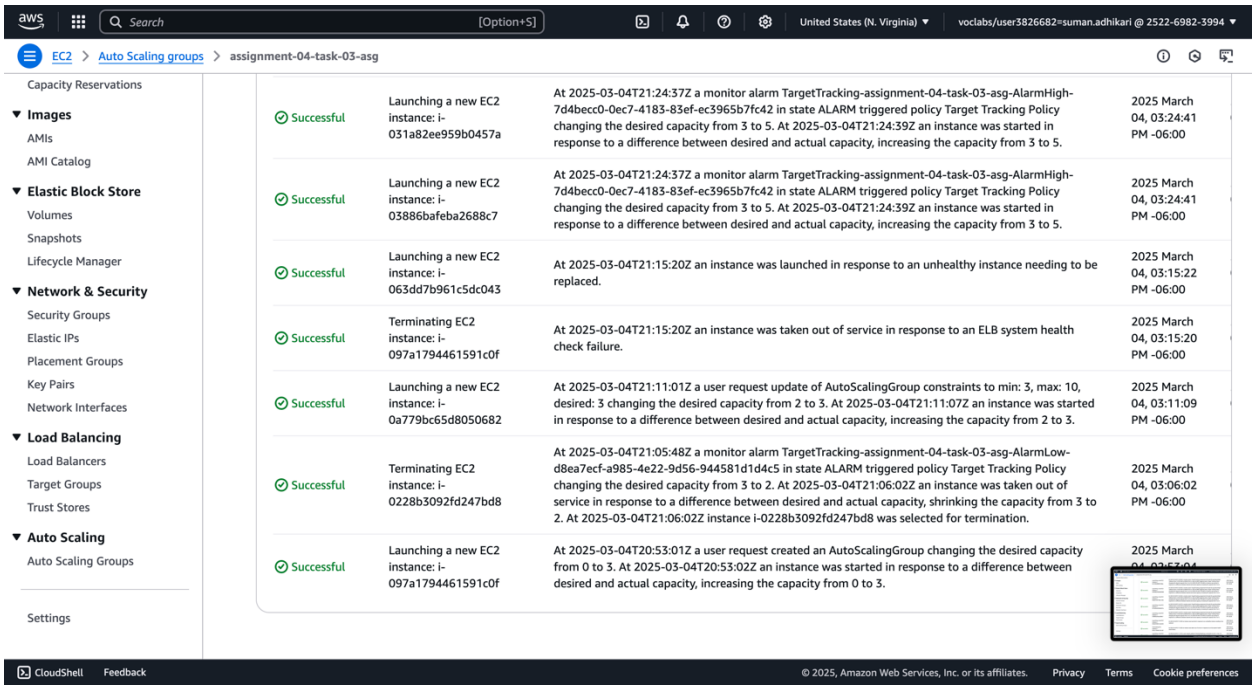
aws
Search
[Option+S]
United States (N. Virginia)
voclabs/user3826682=suman.adhikari @ 2522-6982-3994
EC2 > Auto Scaling groups > assignment-04-task-03-asg
Capacity Reservations
Images
AMIs
AMI Catalog
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
Load Balancing
Load Balancers
Target Groups
Trust Stores
Auto Scaling
Auto Scaling Groups
Settings
Successful
Launching a new EC2 instance: i-031a82ee959b0457a
At 2025-03-04T21:24:37Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmHigh-7d4becc0-0ec7-4183-83ef-ec3965b7fc42 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 3 to 5. At 2025-03-04T21:24:39Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 3 to 5.
2025 March 04, 03:24:41 PM -06:00
Successful
Launching a new EC2 instance: i-03886bafeba2688c7
At 2025-03-04T21:24:37Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmHigh-7d4becc0-0ec7-4183-83ef-ec3965b7fc42 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 3 to 5. At 2025-03-04T21:24:39Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 3 to 5.
2025 March 04, 03:24:41 PM -06:00
Successful
Launching a new EC2 instance: i-063dd7b961c5dc043
At 2025-03-04T21:15:20Z an instance was launched in response to an unhealthy instance needing to be replaced.
2025 March 04, 03:15:22 PM -06:00
Successful
Terminating EC2 instance: i-097a1794461591c0f
At 2025-03-04T21:15:20Z an instance was taken out of service in response to an ELB system health check failure.
2025 March 04, 03:15:20 PM -06:00
Successful
Launching a new EC2 instance: i-0a779bc65d8050682
At 2025-03-04T21:11:01Z a user request update of AutoScalingGroup constraints to min: 3, max: 10, desired: 3 changing the desired capacity from 2 to 3. At 2025-03-04T21:11:07Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 2 to 3.
2025 March 04, 03:11:09 PM -06:00
Successful
Terminating EC2 instance: i-0228b3092fd247bd8
At 2025-03-04T21:05:48Z a monitor alarm TargetTracking-assignment-04-task-03-asg-AlarmLow-d8ea7ecf-a985-4e22-9d56-944581d1d4c5 in state ALARM triggered policy Target Tracking Policy changing the desired capacity from 3 to 2. At 2025-03-04T21:06:02Z an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 3 to 2. At 2025-03-04T21:06:02Z instance i-0228b3092fd247bd8 was selected for termination.
2025 March 04, 03:06:02 PM -06:00
Successful
Launching a new EC2 instance: i-097a1794461591c0f
At 2025-03-04T20:53:01Z a user request created an AutoScalingGroup changing the desired capacity from 0 to 3. At 2025-03-04T20:53:02Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 0 to 3.
2025 March
CloudShell
Feedback
© 2025, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

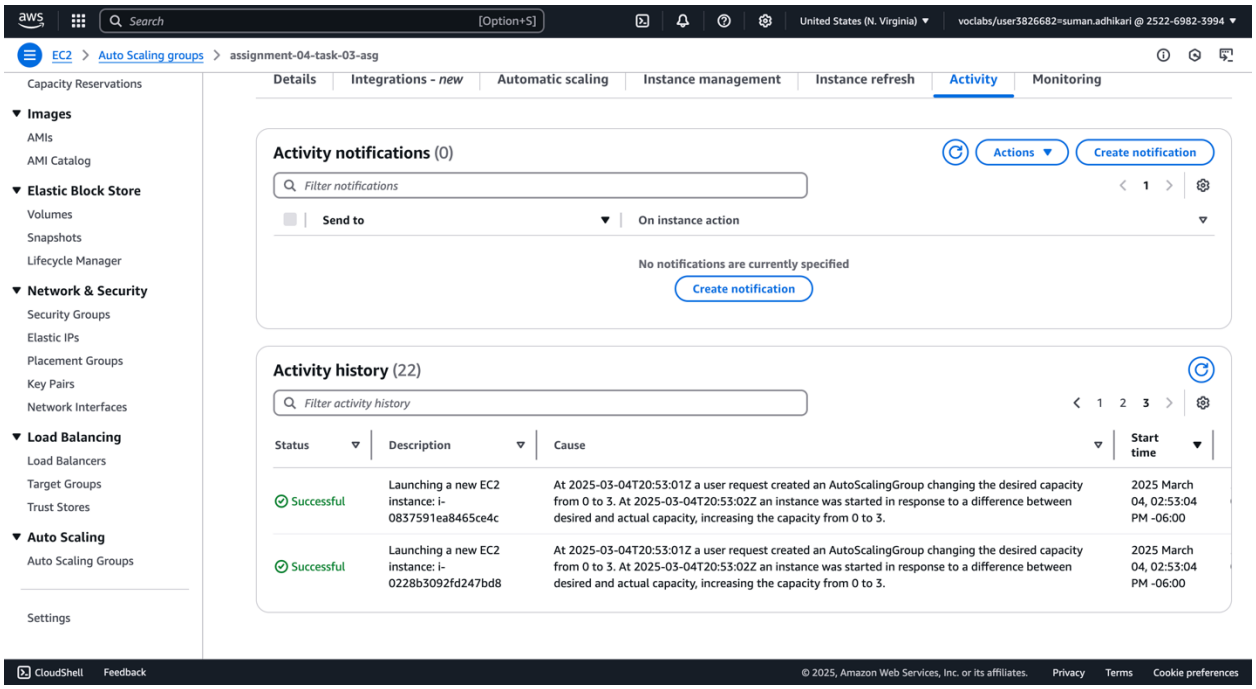

6. Check response from the ALB now.

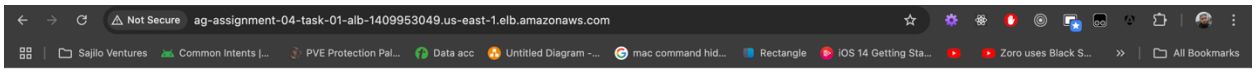

Hello Suman from ip-172-31-84-35.ec2.internal

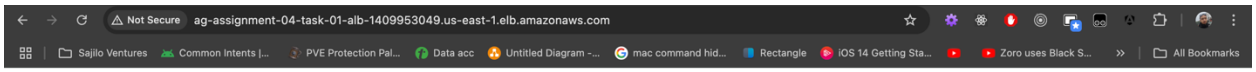

Hello Suman from ip-172-31-35-216.ec2.internal

Task #4 – Cleaning up all services

Nothing much to see here. First, we set the maximum, minimum and desired capacity to 0 in ASG. Then we remove the ASG. Next, the ALB and NLB are deleted. Then we remove then target groups. Afterwards, we terminate any EC2 instances still running on the server. Optionally, remove unused security groups as well.

1. Set capacity to 0 (min, max, desired) in ASG and delete ASG.

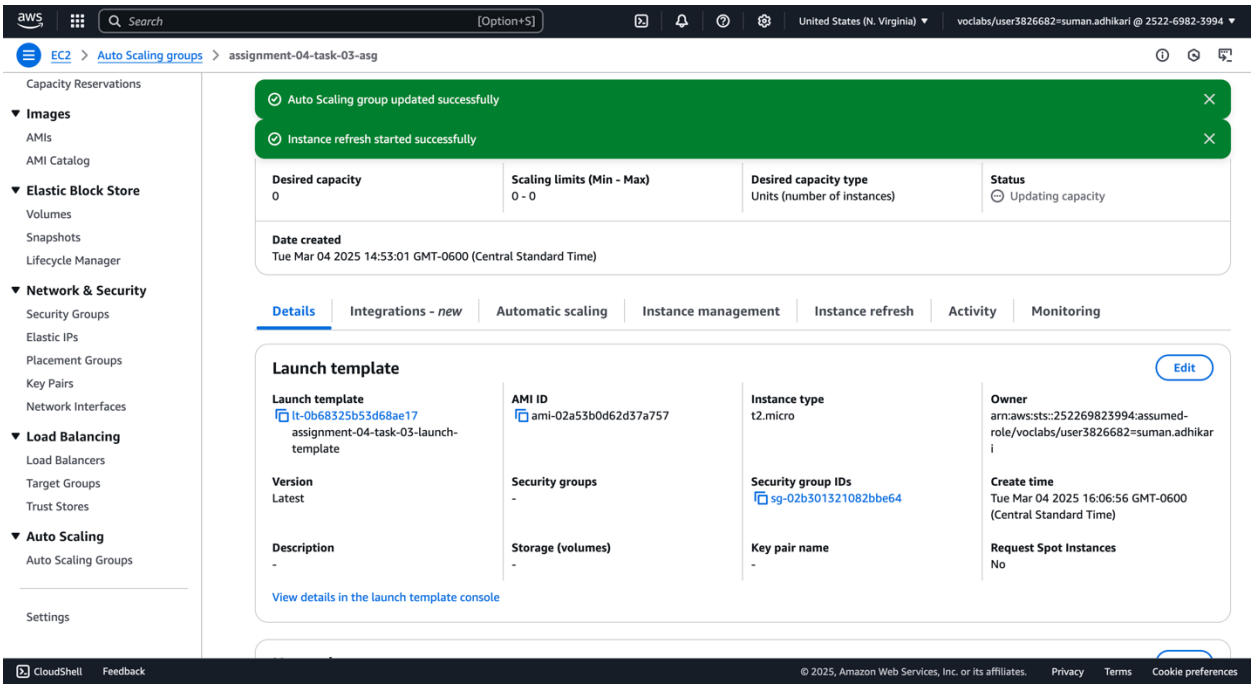

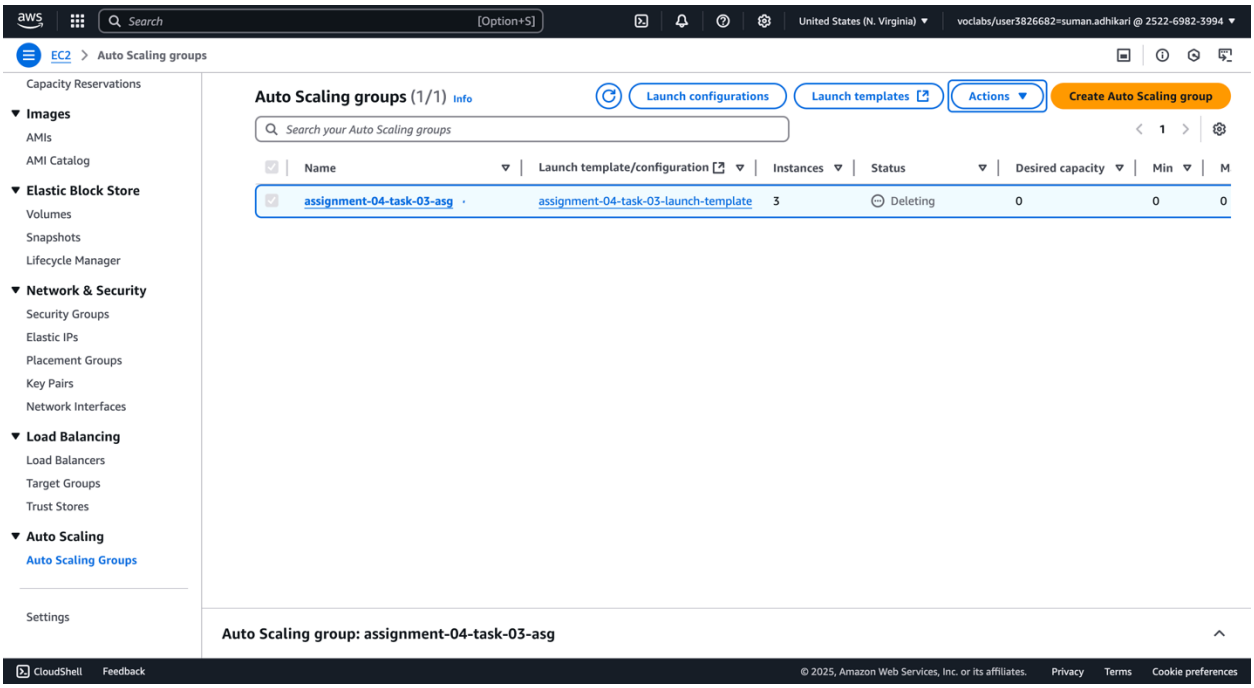

2. Delete ALB and NLB load balancers.

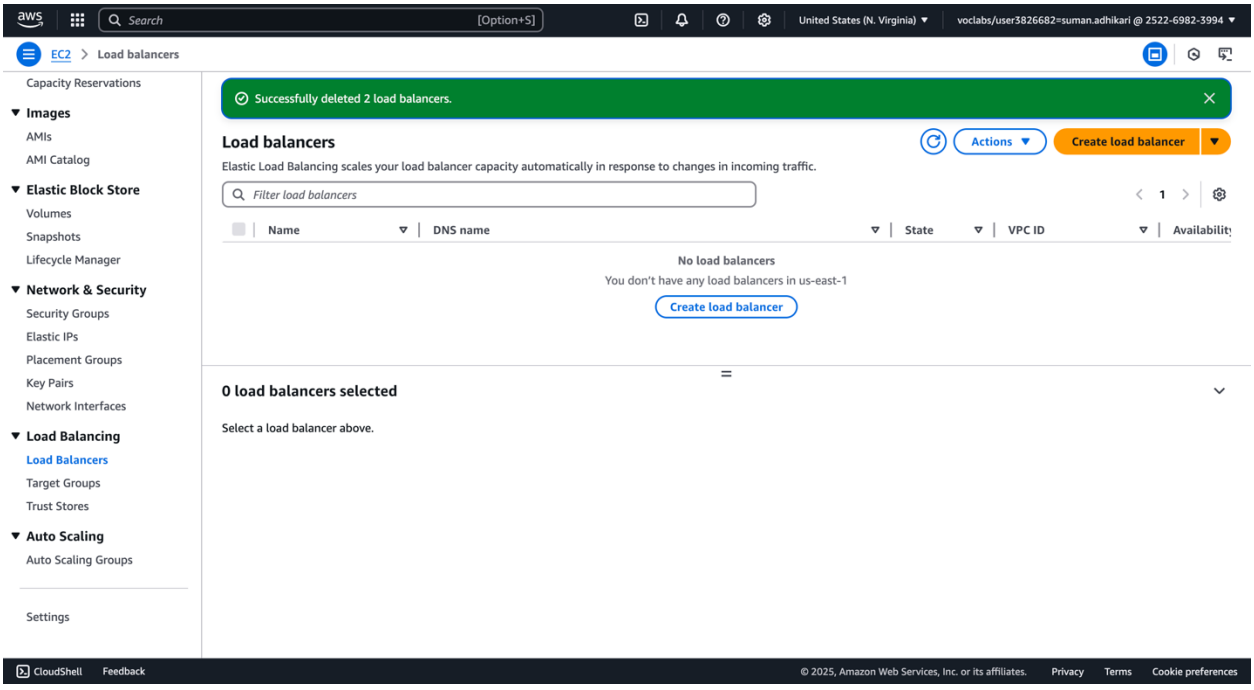

3. Delete target groups.

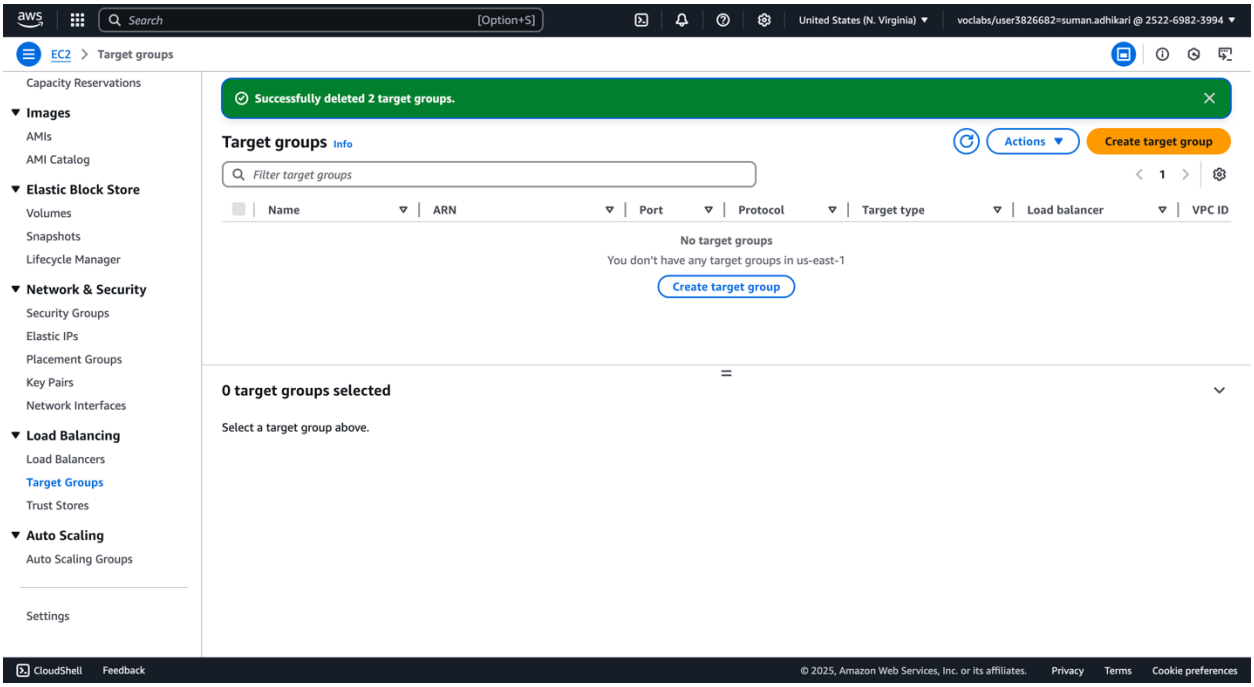

4. Delete Launch Template.

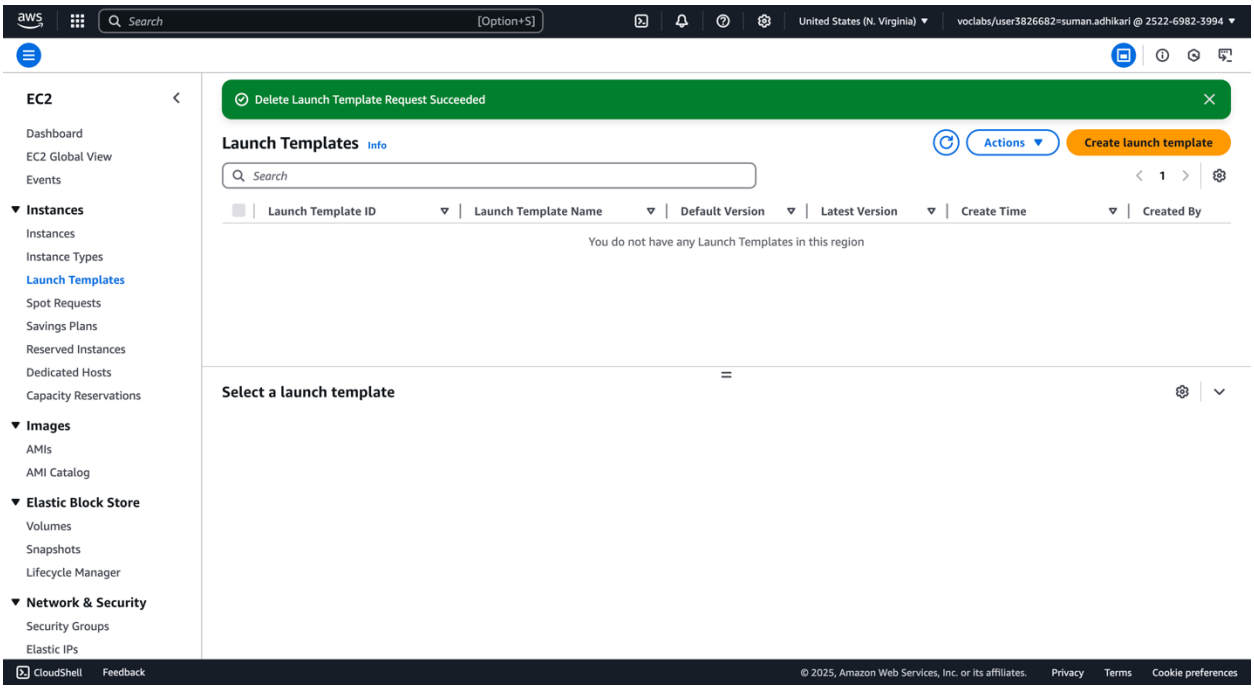

5. Delete any pending instances.